

# Minitutorial: Multi-Modal Data driven and Physics-informed ML + UQ for Materials Applications

## The Hunt for Material Fingerprints



Nat Trask  
University of Pennsylvania  
Sandia National Laboratories



Sandia  
National  
Laboratories



SEA-CROGS



[www.github.com/natrask/PIMA\\_torch](https://www.github.com/natrask/PIMA_torch)

# Talk Overview

## Material Fingerprinting

Our wish list for how to get ML working to integrate physics alongside physics-agnostic experimental modalities

## Variational Inference 101

The basics of applying Bayesian reasoning in the context of deep learning models

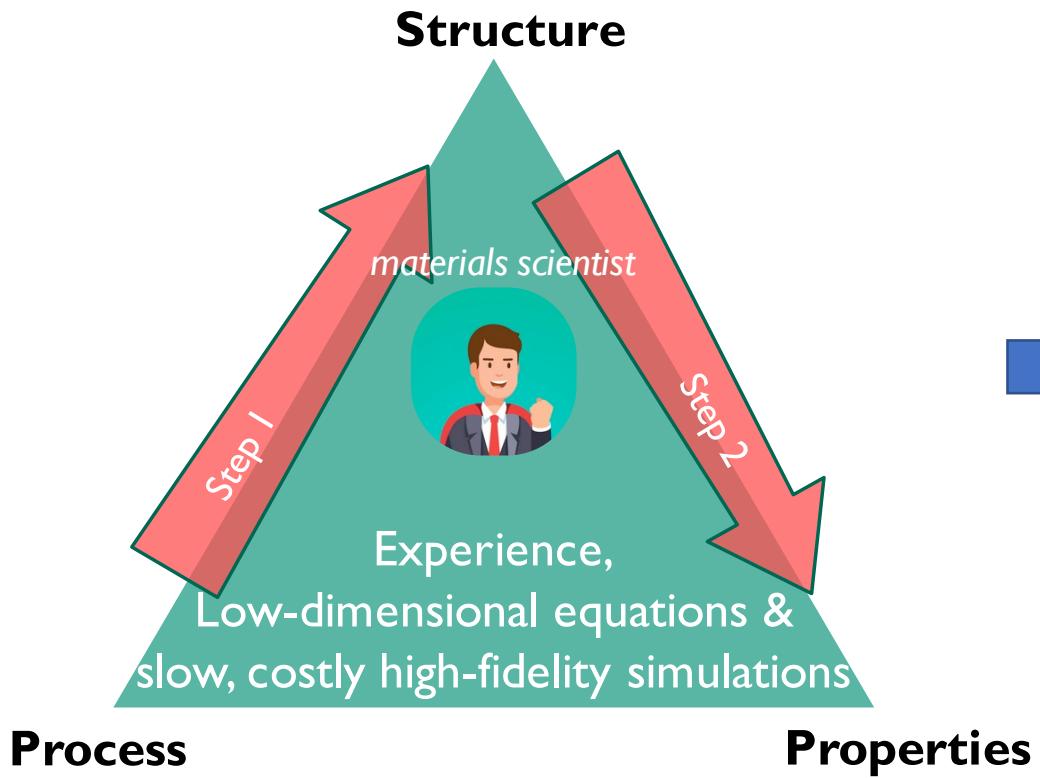
[From ELBO to VAE](#)  
[Clustering with EM](#)

[Physics-based expert models](#)

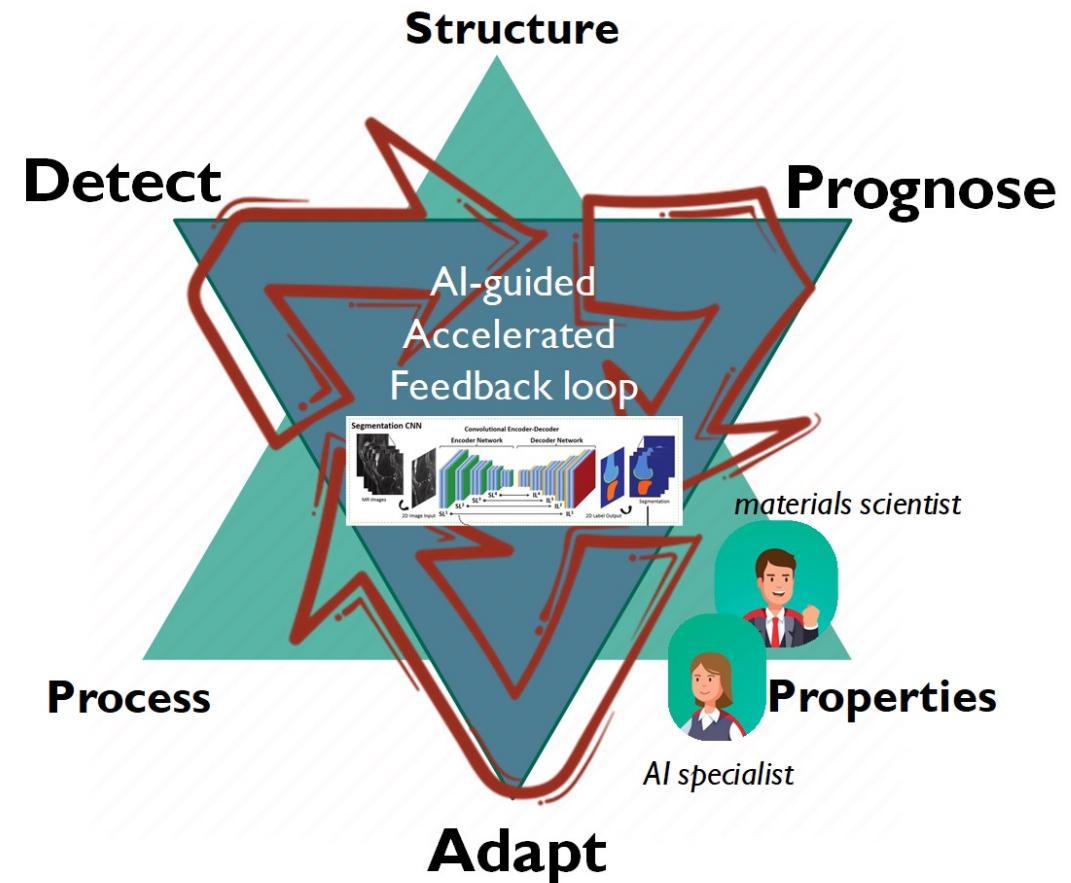
## Physics-Informed Multimodal Autoencoders (PIMA)

A scaled-up production framework for high-throughput materials optimization and discovery

# Beyond bespoke experiments through “self-driving” labs

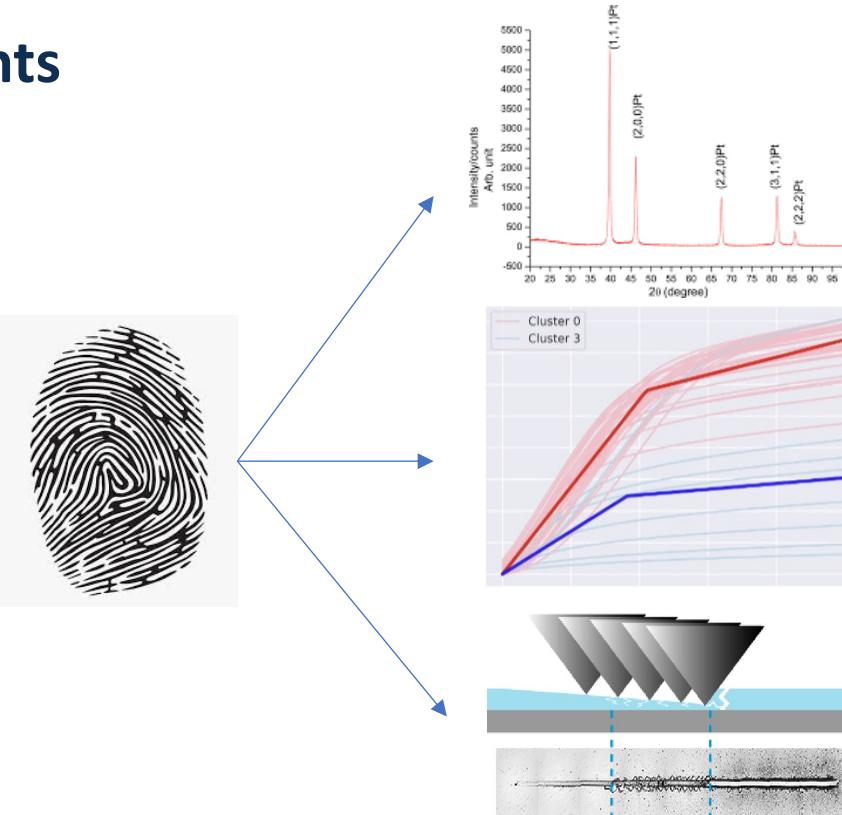
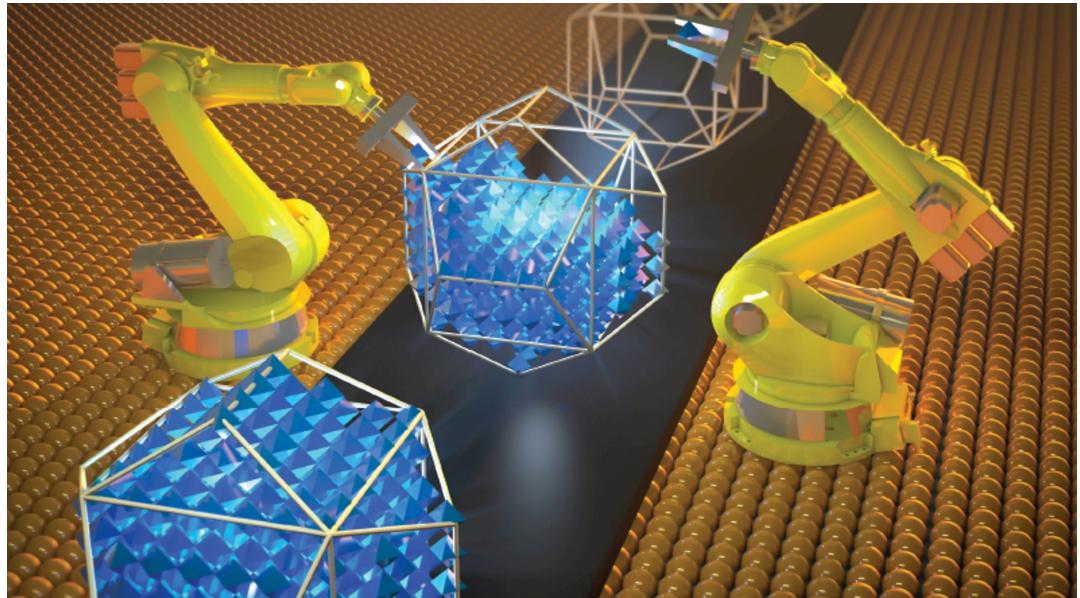


**Traditional material science:**  
**Small** slow data, **trustworthy**  
**Bespoke** single modality experiments  
Campaigns of **1's-10's** of experiments



**High-throughput self-driving labs:**  
**Large** fast data, noisy, **unlabeled**  
Combinatorial, **multimodal** experiments  
Campaigns of **thousands** of experiments

# The hunt for comprehensive material fingerprints



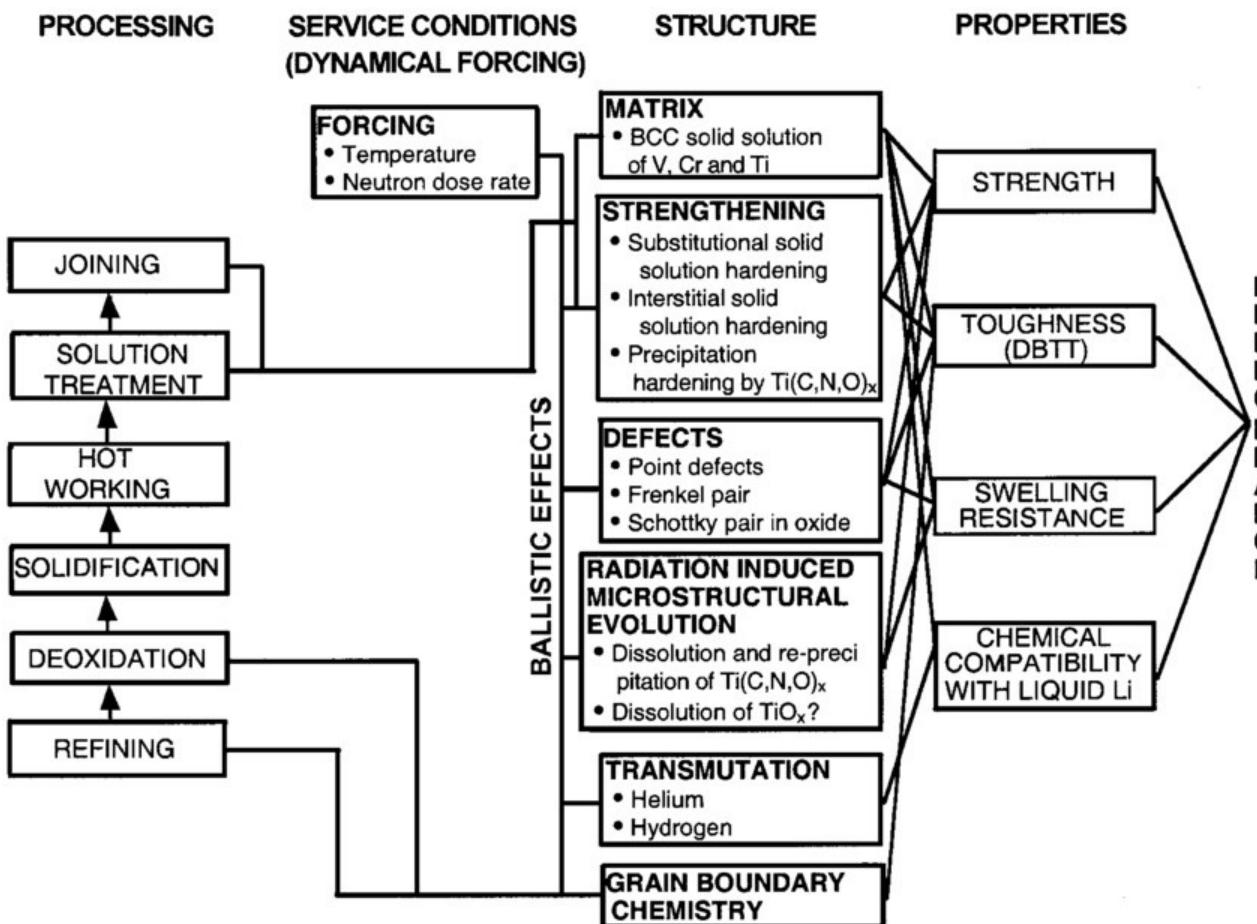
**Fingerprint:** low dimensional signal correlating with an underlying complex physical process in high dimensional data  
**State-of-the-Art:** PCA, feature engineering, manifold learning

**Question: How to move beyond task-agnostic dimension reduction tools to integrate physics?**

Chakraborty, A., Nandi, P., and Chakraborty, B. Fingerprints of the quantum space-time in time-dependent quantum mechanics: An emergent geometric phase. *arXiv preprint arXiv:2110.04370*, 2021.

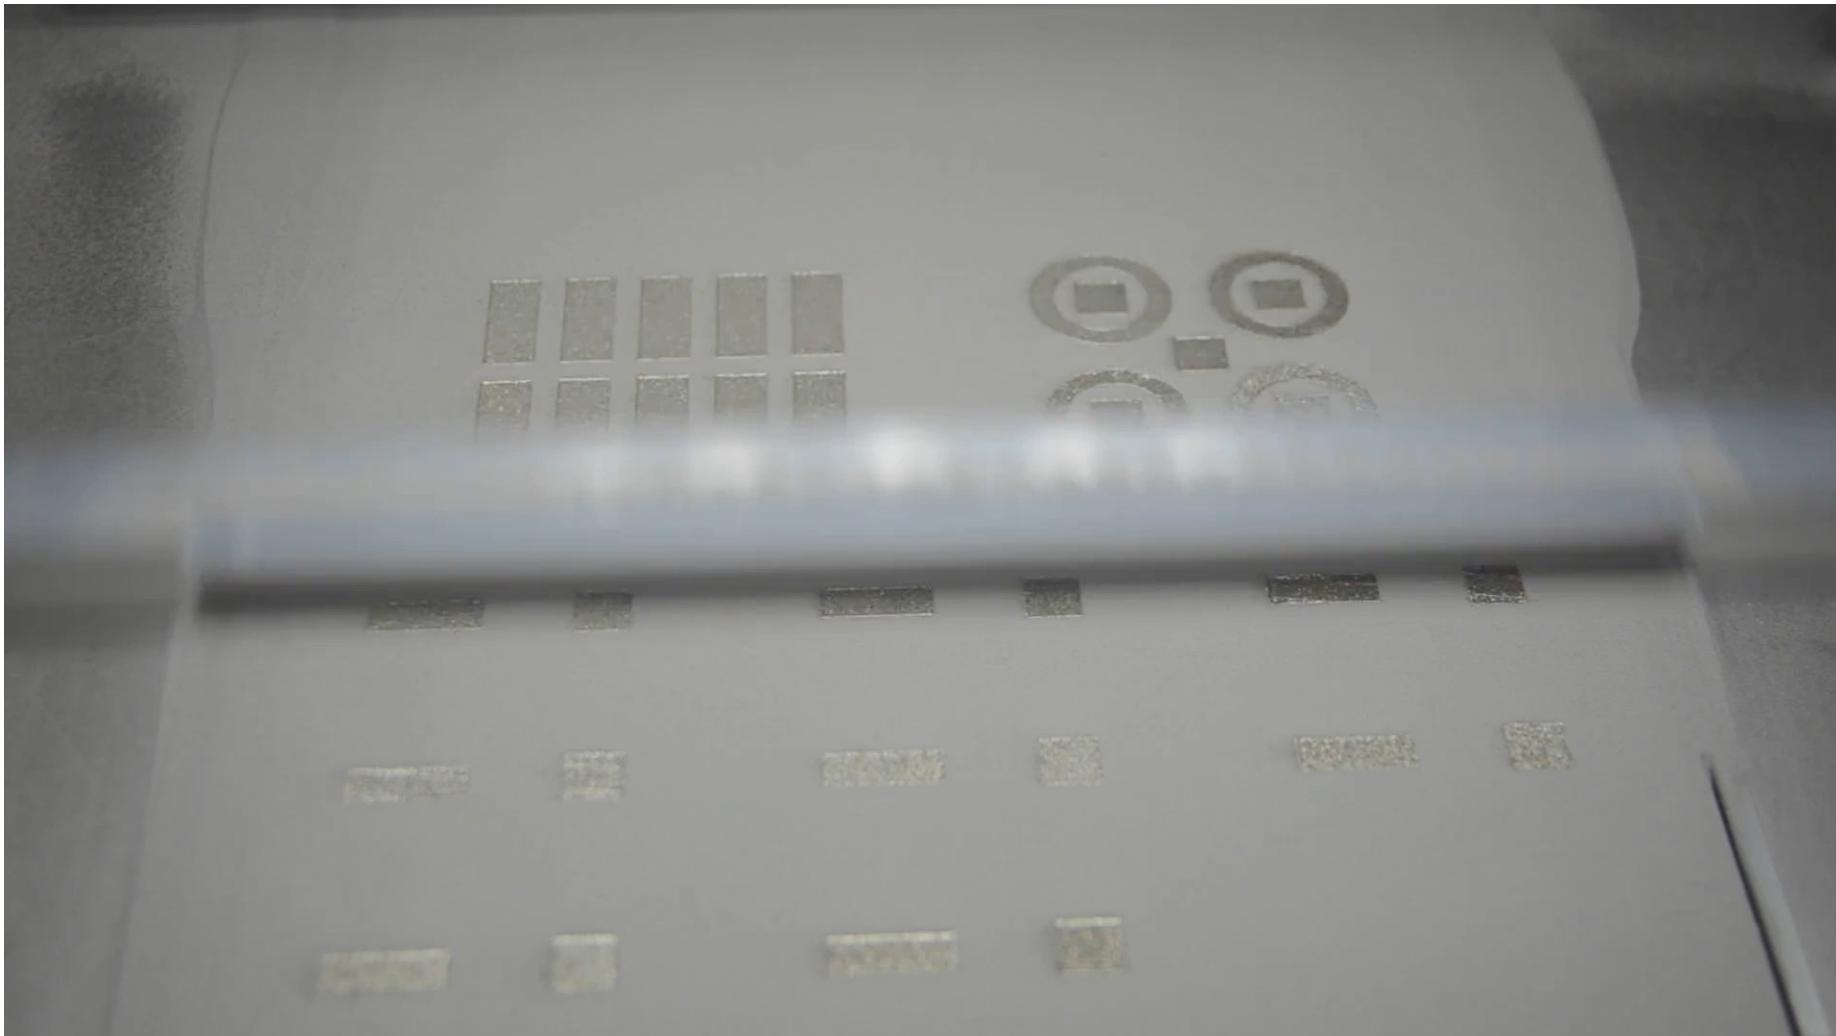
Hasselmann, K. Multi-pattern fingerprint method for detection and attribution of climate change. *Climate dynamics*, 13(9):601–611, 1997.

# Toward process-structure-property-maps with causal & mechanistic learning



When people say AI (as opposed to ML) often they envision  
modular reasoning and cause/effect  
(e.g. Olson Diagram)

## Unconventional sources of multimodal data



**Process parameters:** laser power, speed, path, powder composition

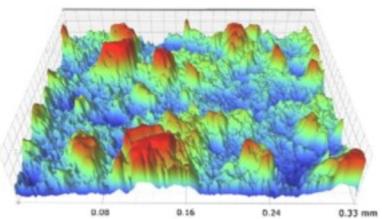
**High-fidelity characterization:** microscopy, XRF, XRD, TEM, SEM

**Low-fidelity signals:** light, sound, images, profilometry

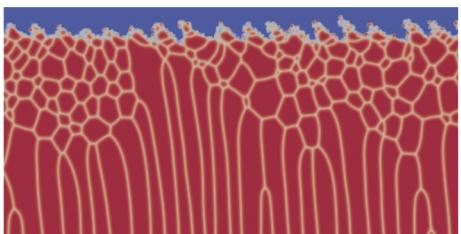
# The avalanche of modalities available in material science applications

## Pre-process

Stage	Composition	Concentration
Electroless nickel plating	$\text{NiSO}_4 \cdot 6\text{H}_2\text{O}$ $\text{NaH}_2\text{PO}_2 \cdot \text{H}_2\text{O}$ $\text{Na}_3\text{C}_6\text{H}_5\text{O}_7 \cdot 2\text{H}_2\text{O}$ $\text{Ni}(\text{SO}_4\text{NH}_3)_2 \cdot 4\text{H}_2\text{O}$ $\text{NiCl}_2 \cdot 6\text{H}_2\text{O}$ $\text{H}_3\text{BO}_3$	$30 \text{ g L}^{-1}$ $30 \text{ g L}^{-1}$ $20 \text{ g L}^{-1}$ $300 \text{ g L}^{-1}$ $15 \text{ g L}^{-1}$ $20 \text{ g L}^{-1}$
Electroplating		

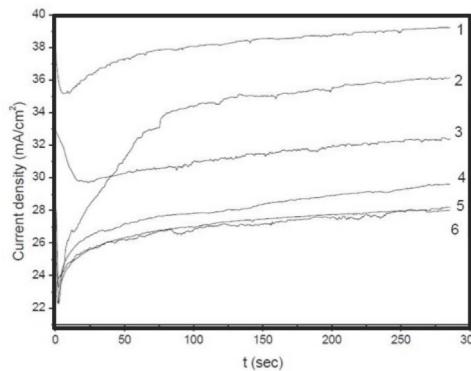


Precursor characterization

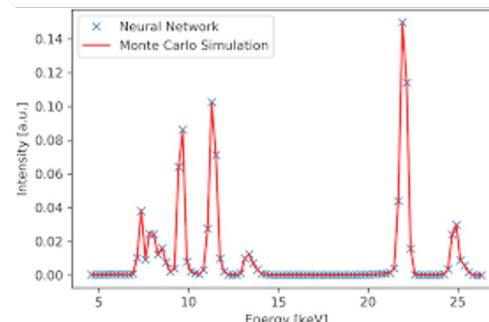


High-fidelity mod/sim

## In-process

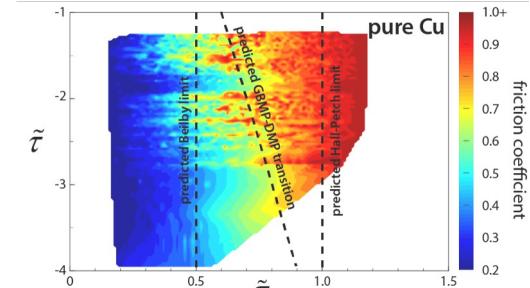
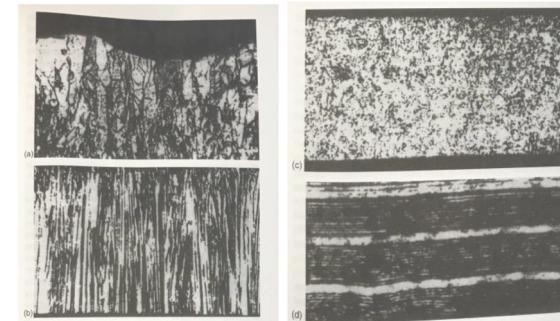


Time-series data



In-situ characterization

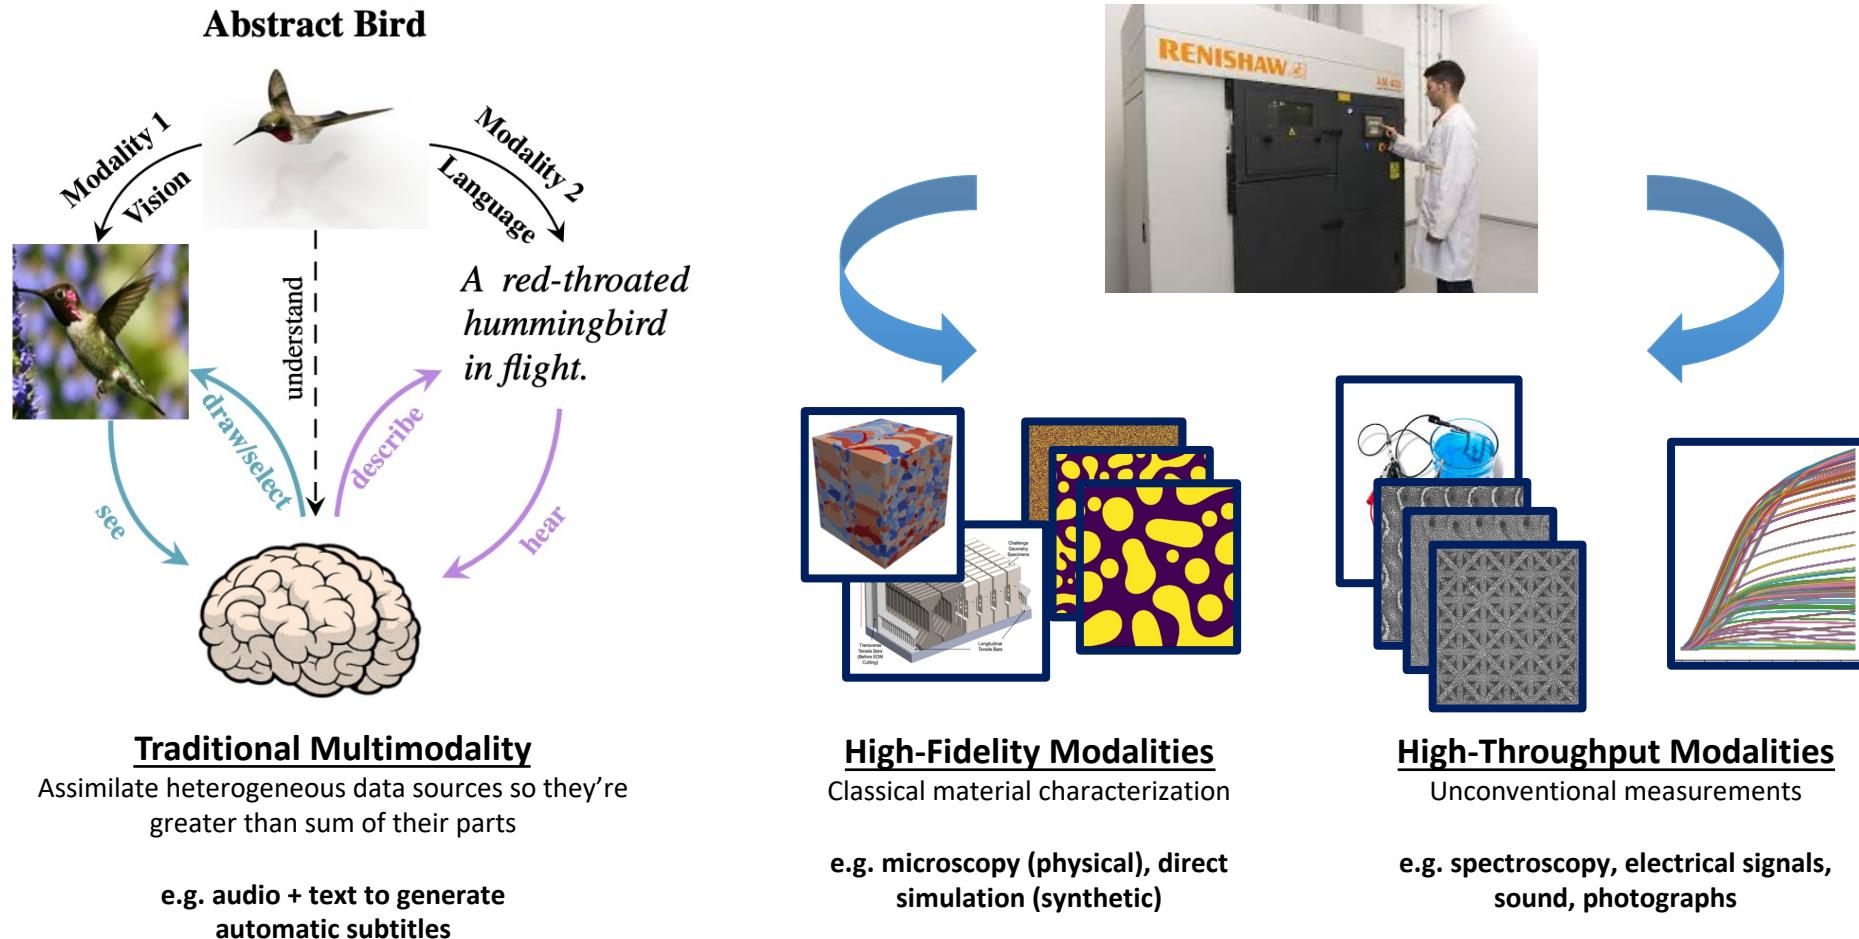
## Post-process



Property/Performance/Aging Measurements

Sifting through a quantity of modalities beyond limits of human cognition

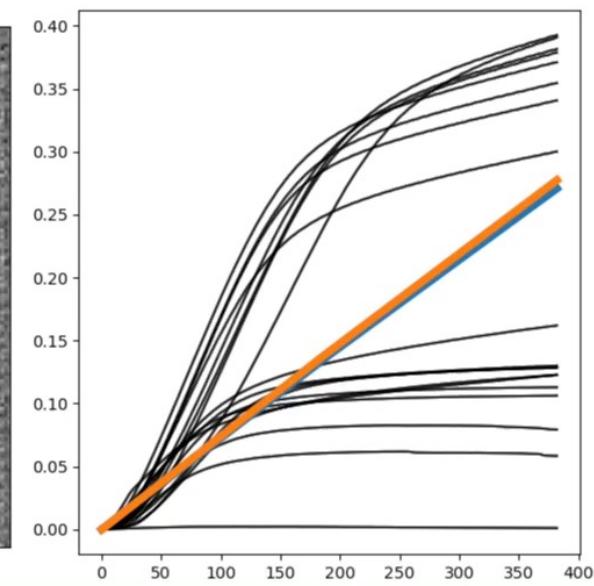
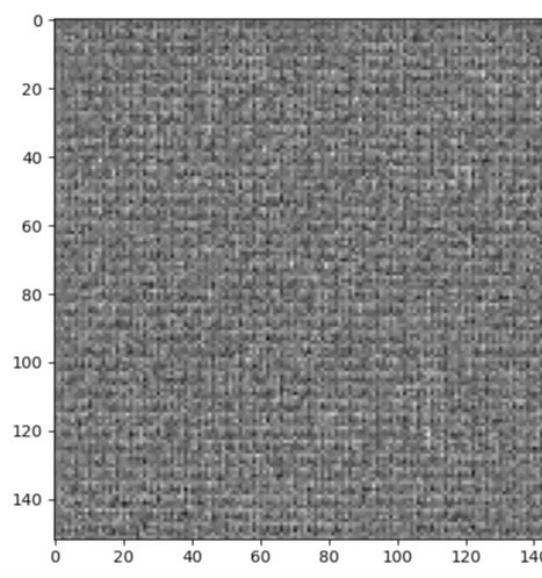
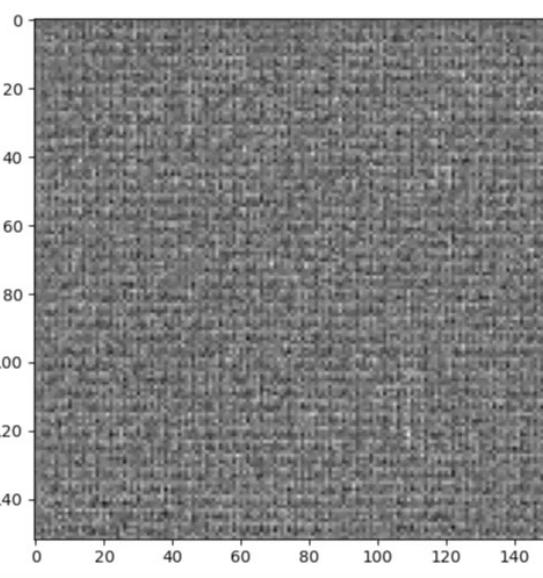
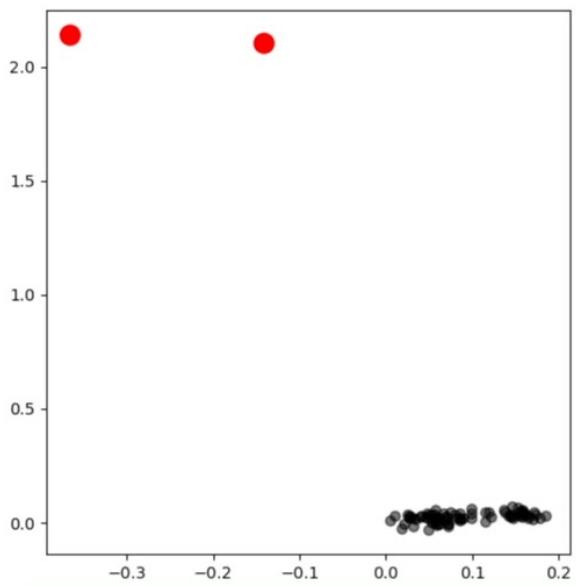
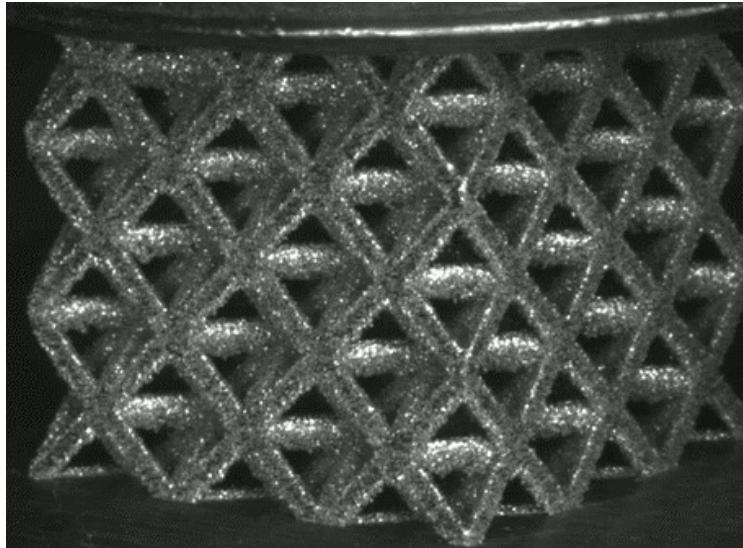
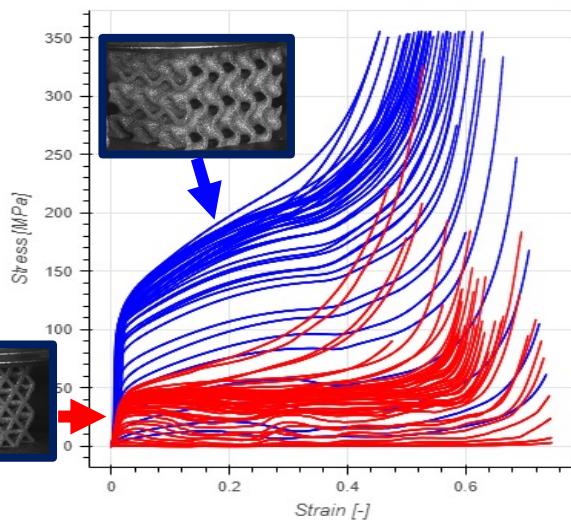
# What do we mean by **multimodal** deep learning?



Shi, Yuge, Brooks Paige, and Philip Torr.  
"Variational mixture-of-experts autoencoders  
for multi-modal deep generative  
models." *Advances in neural information  
processing systems* 32 (2019).

**AI4SS + AI@DOE workshops denote multimodal scientific data a top opportunity for scientific impact across DOE in next 5-10 years**

# Building toward: Disentangling latent representations through physics



1. Images of lattices and stress/strain curves are disentangled in latent space
2. Once disentangled, images are calibrated to clusters
3. Once disentangled, expert models calibrate into two populations
4. **Impact:** Low-throughput uniaxial testing replaced w/ high-throughput imaging

# Talk Overview

## Material Fingerprinting

Our wish list for how to get ML working to integrate physics alongside physics-agnostic experimental modalities

## Variational Inference 101

The basics of applying Bayesian reasoning in the context of deep learning models

[From ELBO to VAE](#)

[Clustering with EM](#)

[Physics-based expert models](#)

## Physics-Informed Multimodal Autoencoders (PIMA)

A scaled-up production framework for high-throughput materials optimization and discovery

# Tutorial 1: maximum likelihood estimates for supervised learning

$$\begin{aligned}\mathcal{L} &= \log p(\mathbf{y}; \theta) = \log \prod_d p(y_d; \theta) \\ &= \sum_d \log p(y_d; \theta) \\ &= \sum_d \log \mathcal{N}(y_d; \mu, \Sigma)\end{aligned}$$

Define loss function as assumed IID joint distribution

Select distribution for marginal distributions

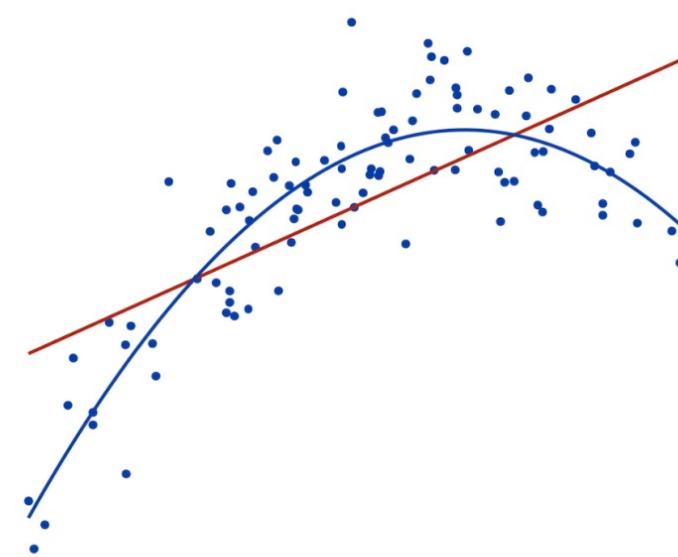
Solve for distribution parameters which maximize  
(log) likelihood of your dataset

**Exercise 1:** Recast polynomial fitting as an MLE problem

$$\Theta = \left\{ \mu(x) = \sum_{\alpha} c_{\alpha} \phi_{\alpha}(x), \quad \sigma^2 \right\}$$

$$p(y_d | \Theta) = \mathcal{N}(\mu(x_d), \sigma)$$

Show:  $\mathcal{L} = C + N \log |\sigma| + \sum_{d=1}^N \left( \frac{y_d - \sum_{\alpha} c_{\alpha} \phi_{\alpha}(x_d)}{\sigma} \right)^2$



# MLE is intractable for **hidden/latent** data

**Given:**

Hidden variables -  $Z$

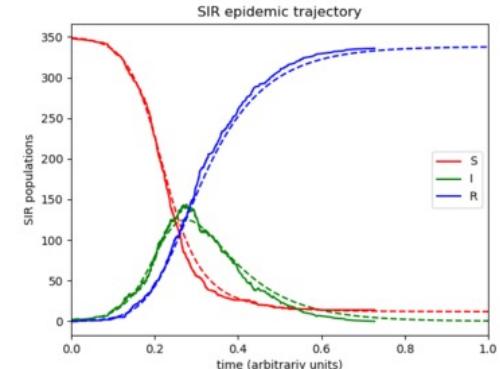
Observable variables -  $X$

$$\mathcal{L} = - \sum_d \log p(X_d, Z_d)$$

Can't plug data in for  $Z$ !

$$\begin{cases} \frac{dS}{dt} = -\beta IS, \\ \frac{dI}{dt} = \beta IS - \gamma I, \\ \frac{dR}{dt} = \gamma I, \end{cases}$$

where  $S$  is the stock of susceptible population,  $I$  is the stock of infected,  $R$  is the stock of removed population (either by death or recovery), and  $N$  is the sum of these three.



Example: For compartment models during COVID,  
no direct measurements of  $I$  were available

Can we estimate the posterior distribution:

$$p(Z|X) = \frac{p(X|Z)p(Z)}{p(X)} = \frac{p(X|Z)p(Z)}{\underbrace{\int p(X|Z)p(Z)dZ}_{\text{Computationally intractable integral}}}$$

**Option 1: (Sampling)**

Use data intensive sampling to estimate denominator – can't afford!

**Option 2: (Variational Inference)**

Solve an optimization problem to infer a “close as possible” posterior distribution that’s “nice” to work with

**Variational inference 101:** First, replace log likelihood with something that doesn't depend on  $z$

$$\mathcal{L} = - \sum_d \log p(y_d | \Theta)$$

Start with what we know

$$= - \sum_d \log \sum_{z_d} p(y_d, z_d | \Theta)$$

Marginalization / LoTP

$$= - \sum_d \log \sum_{z_d} q(z_d) \frac{p(y_d, z_d | \Theta)}{q(z_d)}$$

Introduce an arbitrary distribution  $q$  and multiply by 1

$$= - \sum_d \log \mathbb{E}_q \left[ \frac{p(y_d, z_d | \Theta)}{q(z_d)} \right]$$

Reinterpret inner sum as an expectation with respect to the measure  $q$

$$\leq - \sum_d \mathbb{E}_q \left[ \log \frac{p(y_d, z_d | \Theta)}{q(z_d)} \right]$$

Apply Jensen's inequality

$$= - \sum_d \mathbb{E}_q [\log p(y_d, z_d | \Theta)] - H(q_d)$$

Entropy of  $q$

**Evidence Lower BOund (ELBO)**

Lower bound on likelihood considering all possible  $q$ 's



# How do we make the ELBO as tight as possible of a lower bound on the LL?

$$\begin{aligned}
 \epsilon(\Theta, q_d | y_d) &= \sum_d \mathbb{E}_q \left[ \log \frac{p(y_d, z_d | \Theta)}{q_d} \right] \\
 &= \sum_d \mathbb{E}_q \left[ \log \frac{p(z_d | y_d, \Theta) p(y_d | \Theta)}{q_d} \right] \\
 &= \sum_d \mathbb{E}_q \left[ \log \frac{p(z_d | y_d, \Theta)}{q_d} \right] + \mathbb{E}_q [\log p(y_d | \Theta)] \\
 &= \sum_d -\text{KL}(p(z_d | y_d, \Theta) || q_d) + \mathbb{E}_q [\log p(y_d | \Theta)]
 \end{aligned}$$

Get a Kullback-Leibler divergence  
 (a quasi-metric on distributions)

**Choose the q that makes ELBO as tight as possible**

Since KL is non-neg, ELBO is biggest when

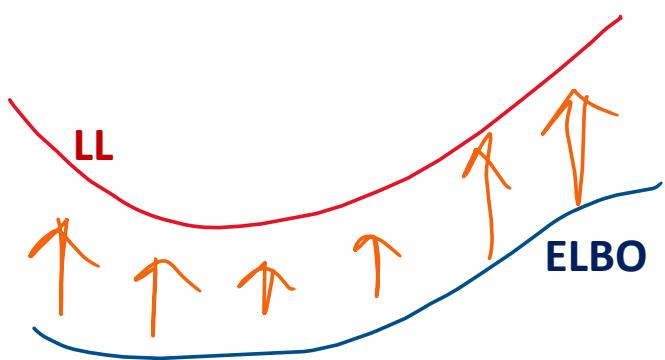
$$q_d = p(z_d | y_d, \Theta)$$

**Evidence Lower Bound (ELBo)**

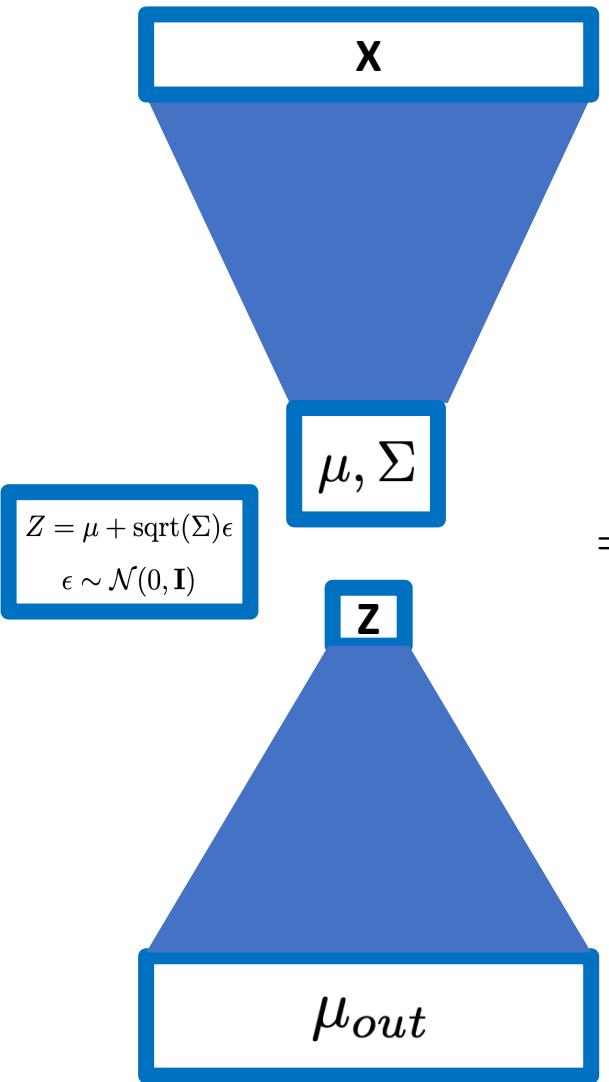
Lower bound on likelihood  
considering all possible q's

Law of total probability to  
separate marginal on y

Linear separability of logs of  
products give 2 terms



# Now that we have an alternative loss, how can we use it?



$$p(Z) \sim \mathcal{N}(0, \mathbf{I})$$

$$q(Z|X) \sim \mathcal{N}(\mu, \Sigma)$$

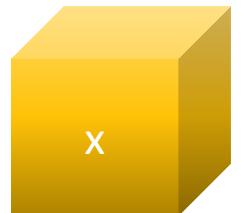
$$p(X|Z) \sim \mathcal{N}(\mu_{out}, \mathbf{I})$$

$$\text{ELBO}(q) = \sum_Z q(Z) \log \frac{p(Z, X)}{q(Z)}$$

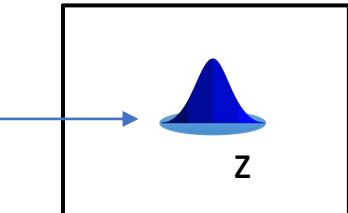
$$= \sum_Z \underbrace{\mathbb{E}_q[\log p(X|Z)]}_{\text{Reconstruction Error (MSE)}} - \underbrace{KL(q(Z|X) || p(Z))}_{\text{Prior/Posterior Mismatch Penalty}}$$

Reconstruction Error (MSE)

Prior/Posterior Mismatch Penalty



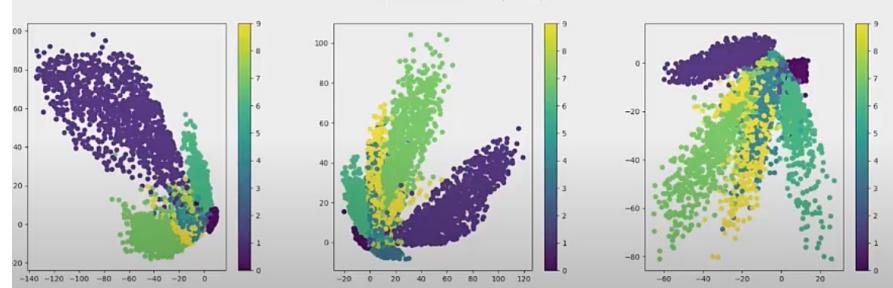
High-dim Data Space



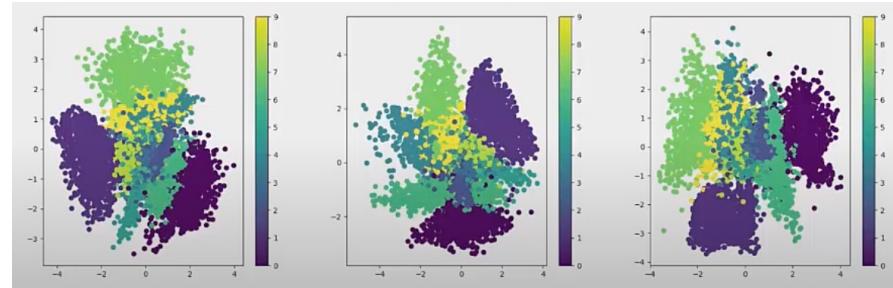
Low-dim Latent Space

[Link to colab notebook](#)

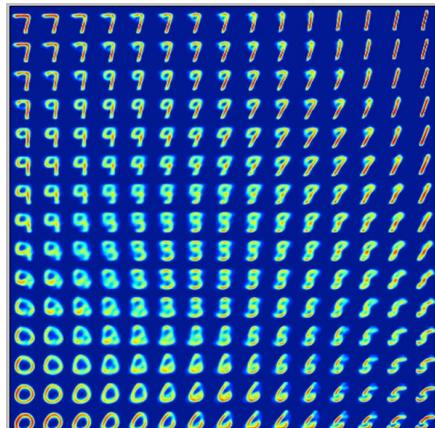
Vanilla Autoencoder



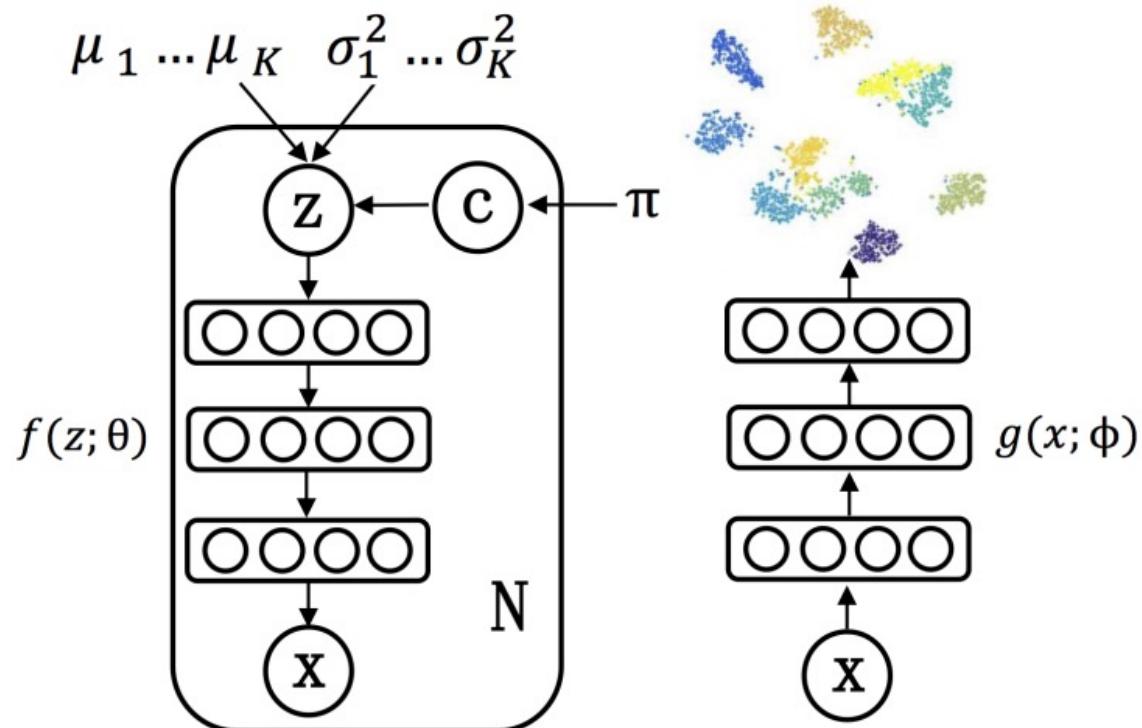
Variational Autoencoder (VAE)



Choose prior to allow generative modeling



There is a lot more we can do with the prior than make a plain VAE though



To obtain manipulable clusters, take GMM prior:

$$p(Z) = \sum_c \Pi_c \mathcal{N}(Z; \mu_c, \Sigma_c)$$

Jiang, Zhuxi, et al. "Variational deep embedding: An unsupervised and generative approach to clustering." *arXiv preprint arXiv:1611.05148* (2016).

Dilokthanakul, Nat, et al. "Deep unsupervised clustering with gaussian mixture variational autoencoders." *arXiv preprint arXiv:1611.02648* (2016).

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5
6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2
8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8
9	4	4	4	4	9	4	4	4	4	4	4	4	4	4	4
7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7
3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3
1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
4	9	4	9	4	9	4	9	4	9	4	9	4	9	4	9

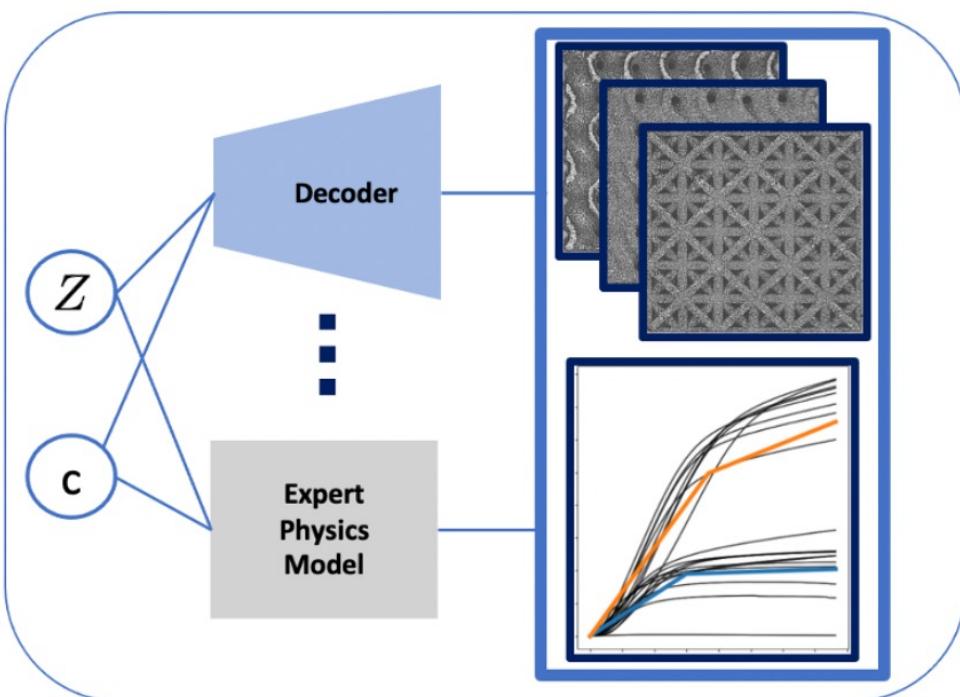
Explicitly access  
disentangled clusters  
describing semantically  
distinct clusters in dataset

Can we steer unsupervised  
clustering algorithms that  
are biased toward  
**physically-biased** clusters  
encoding material  
fingerprints?

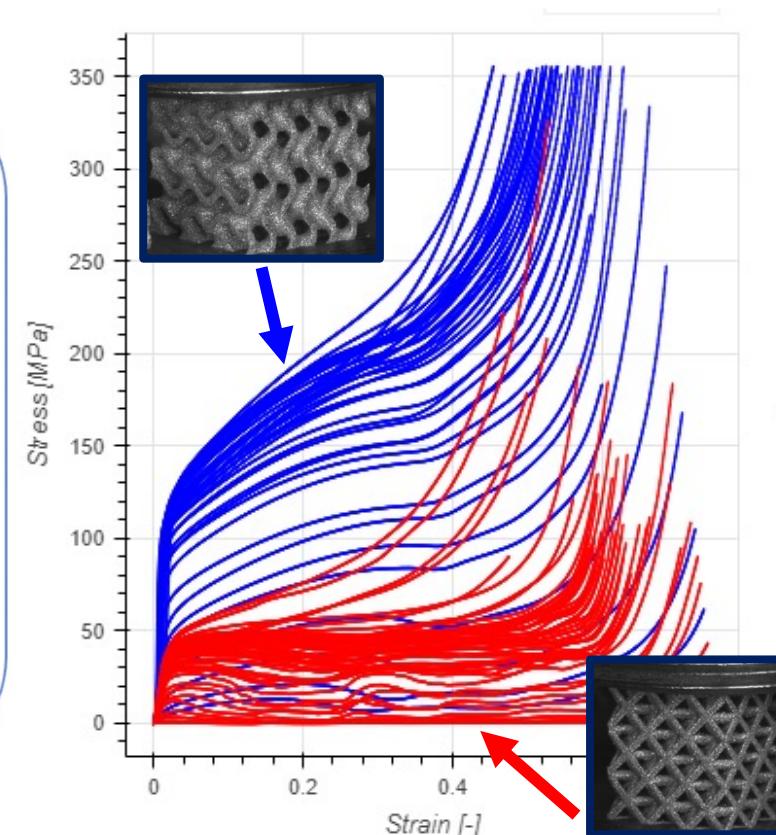
## Tutorial 2: Mixture of expert models conditioned on Gaussian mixtures

$$z \sim GMM \left( \pi, \left\{ \mathcal{N}(\mu_c, \sigma_c^2) \right\}_{c=1}^N \right) \quad c \sim cat(\pi) \quad y|c, t \sim \mathcal{N}(\mu_c = f(t; \Theta_c), \sigma_c^2)$$

Prior Distribution on Latent Space

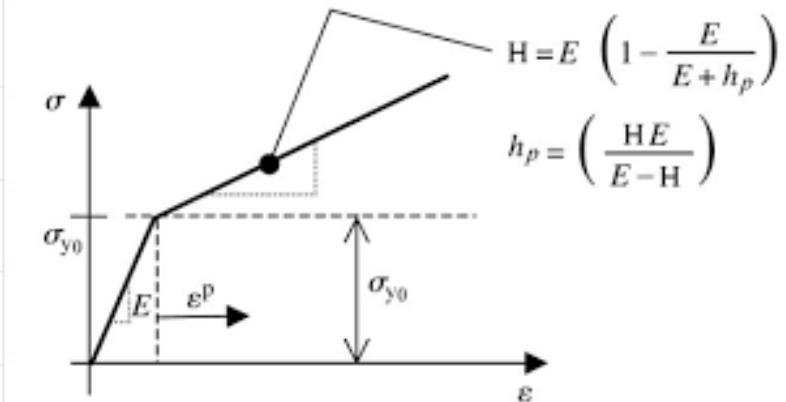


Latent Cluster Probability



Cluster-specific Physics Model

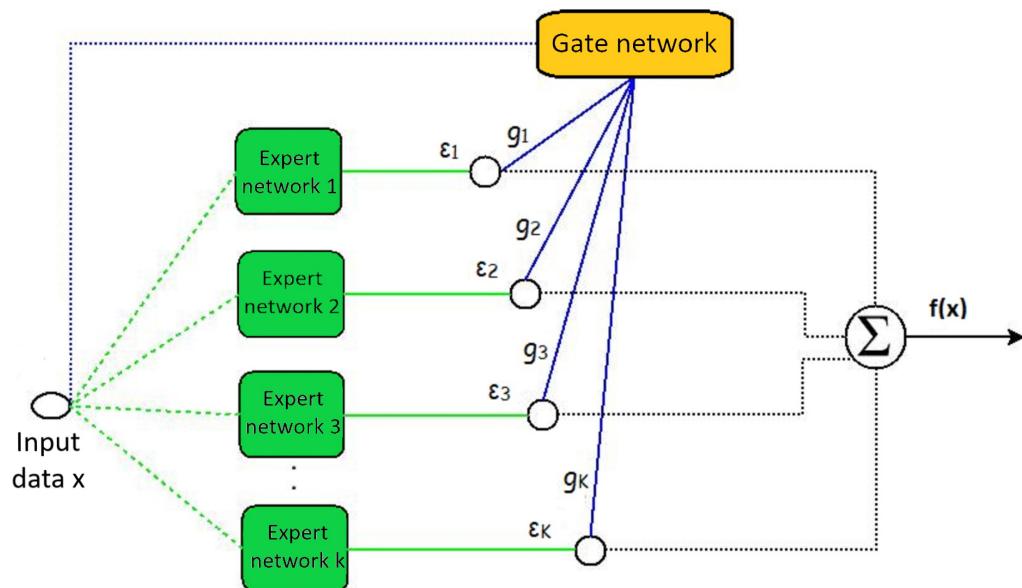
Each cluster can get its own custom tailored physics model



## Tutorial 2: Mixture of expert models conditioned on Gaussian mixtures (simplified)

$$c \sim \text{cat}(\pi) \quad y|c, t \sim \mathcal{N}(\mu_c = f(t; \Theta_c), \sigma_c^2) \quad p(y; \pi, \Theta) = \sum_c \pi_c \mathcal{N}(f(t; \Theta_c), \sigma_c^2)$$

Latent Cluster Probability      Cluster-specific Physics Model      Mixture of Experts Model



Architecture for a contemporary MOE model

### Mixture of Experts (MoE)

Blending expert models is an old idea getting a resurgence recently, as a powerful tool to blend smaller models and access very, very large models [e.g. 100B+ param (*Shazeer 2018*)]

### As an exercise, we use VI to fix constant gates

In the process we will learn *expectation maximization (EM)*, a nice technique for problems with latent variables

Jordan, Michael I., and Robert A. Jacobs. "Hierarchical mixtures of experts and the EM algorithm." *Neural computation* 6.2 (1994): 181-214.

Peralta, Billy, Ariel Saavedra, Luis Caro, and Alvaro Soto. "Mixture of experts with entropic regularization for data classification." *Entropy* 21, no. 2 (2019): 190.

# Exercise: Derive expectation-maximization algorithm for MoE model

## E-step

Calculate  $q = p(c|y)$

## M-step

Optimize ELBO holding  $q$  from  
M-step constant

$$\begin{aligned} q = p(c|y) &= \frac{p(c)p(y|c)}{\sum_{c'} p(c')p(y|c')} && \text{Bayes rule} \\ &= \frac{\pi_c \mathcal{N}(f(t; \theta_c), \sigma_c^2)}{\sum_{c'} \mathcal{N}(f(t; \theta_{c'}), \sigma_{c'}^2)} && \text{Estimate } q \end{aligned}$$

## E-step

Using the current guesses for the model, we compute the posterior distribution  $q$  to identify which model is most likely to have generated the data

[Sample code and blog for EM algorithm](#)

Murphy, Kevin P. *Probabilistic machine learning: an introduction*. MIT press, 2022.

# Exercise: Derive expectation-maximization algorithm for MoE model

**E-step**  
Calculate  $q = p(c|y)$

**M-step**  
Optimize ELBO holding  $q$  from  
M-step constant

$$\min \sum_d \mathbb{E}_q [\log p(y_d, z_d | \Theta)]$$

Expand definition of ELBO

$$= \sum_{d,c} q_{dc} [\log p(y_d | z_d, \Theta_c) + \log p(z_d)]$$

LoTP

$$= \sum_{d,c} q_{dc} \left[ \log \mathcal{N}(f(t; \Theta_c), \sigma_c^2) + \log \pi_c \right]$$

Defs from MoE

$$\frac{\partial}{\partial \pi_c} = 0 \rightarrow \pi_c = \frac{\sum_d q_{dc}}{N}$$

$$\frac{\partial}{\partial \Theta_c} = 0 \rightarrow \theta_c = \operatorname{argmin} \sum_d q_{dc} (y_d - f(t; \Theta_c))^2$$

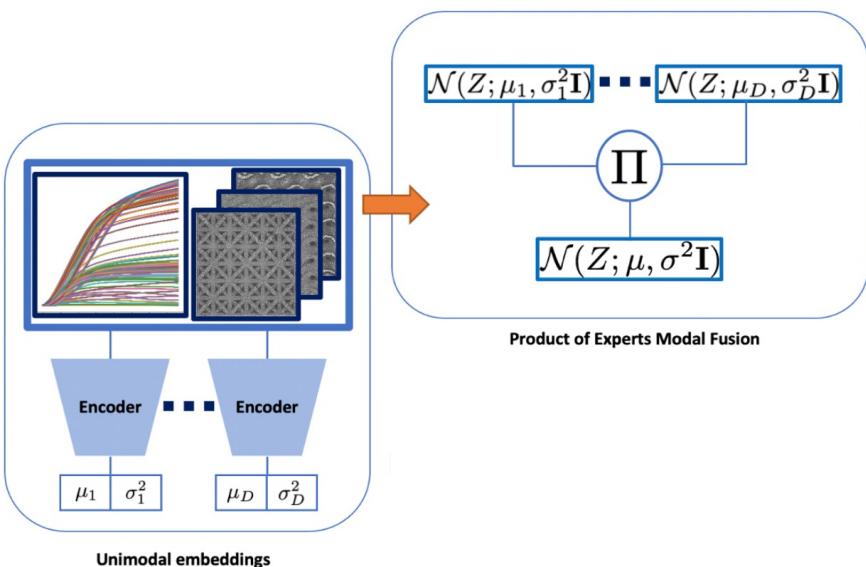
## M-step

Calibration of **all** random variables follows from either explicit formulas or weighted L2 residuals

**Note this for the remainder of the talks in this mini**

[Sample code and blog for EM algorithm](#)

# Final Exercise: Enriching VAE with multimodal product of experts (PoE)



## Key Idea

Use standard VAE encodings, and exploit the fact that product of Gaussian PDFs yields another Gaussian PDF

Derive through  
Bayes + IID  
assumption

$$q(Z|X_1, \dots, X_D) = q(Z)^{1-D} \prod_{i=1}^D q(Z|X_i),$$

Uninformative prior

$$q(Z|X_1, \dots, X_D) \propto \prod_{i=1}^D q(Z|X_i)$$

Closed form posterior

$$q(Z|X_1, \dots, X_D) = \mathcal{N}(\mu, \sigma^2 \mathbf{I}),$$

$$\sigma^{-2} = \sum_{i=1}^D \sigma_i^{-2}, \quad \frac{\mu}{\sigma^2} = \sum_{i=1}^D \frac{\mu_i}{\sigma_i^2},$$

Wu, M., & Goodman, N. (2018). Multimodal generative models for scalable weakly-supervised learning. *Advances in neural information processing systems*, 31.

# Talk Overview

## Material Fingerprinting

Our wish list for how to get ML working to integrate physics alongside physics-agnostic experimental modalities

## Variational Inference 101

The basics of applying Bayesian reasoning in the context of deep learning models

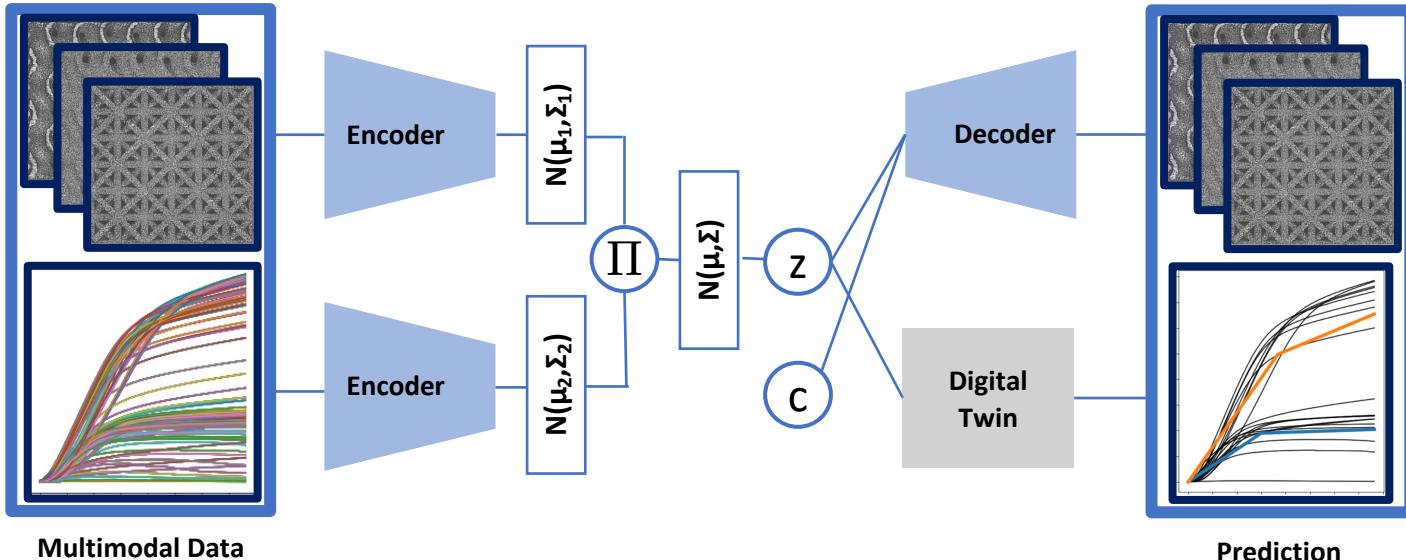
[From ELBO to VAE](#)  
[Clustering with EM](#)

[Physics-based expert models](#)

## Physics-Informed Multimodal Autoencoders (PIMA)

A scaled-up production framework for high-throughput materials optimization and discovery

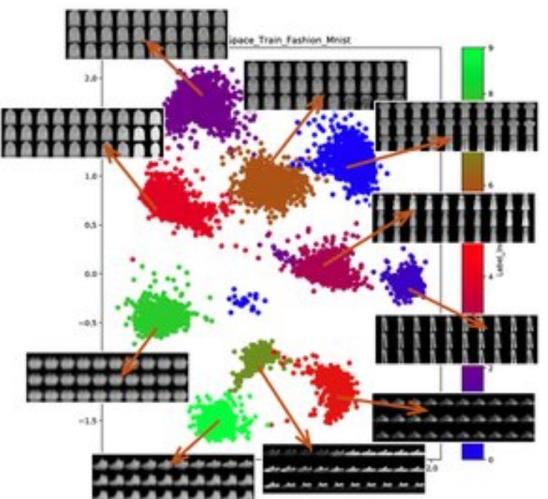
# The complete framework: Physics-informed multimodal autoencoders (PIMA)



Walker, Elise, et al.  
"Unsupervised physics-informed disentanglement of multimodal data."  
*Foundations of Data Science* (2024): 0-0.

Trask, Nathaniel Albert, Carianne Martinez, and Brad Boyce. "Physics-informed multimodal autoencoder." U.S. Patent Application No. 17/743,160.

**Informal Idea:**  
We discover a shared latent representation of data providing a Rosetta stone across modalities w/ uncertainty estimation



**Formal Idea:**

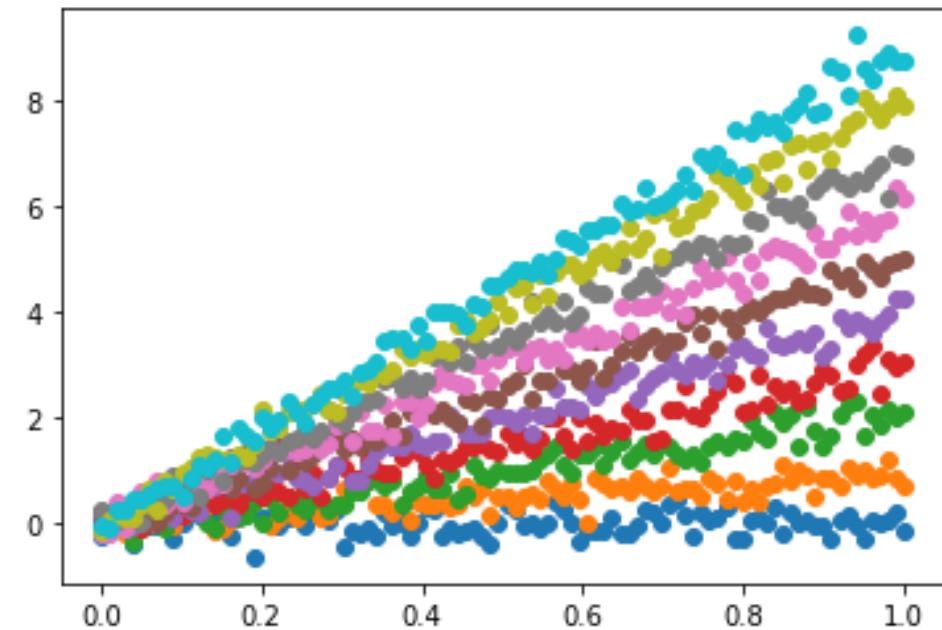
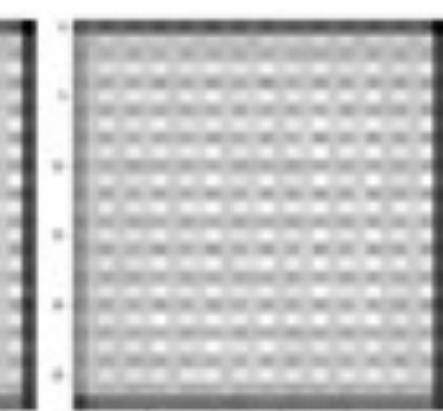
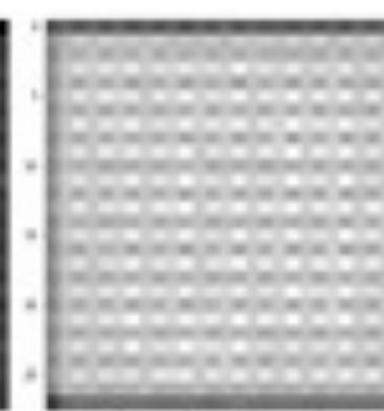
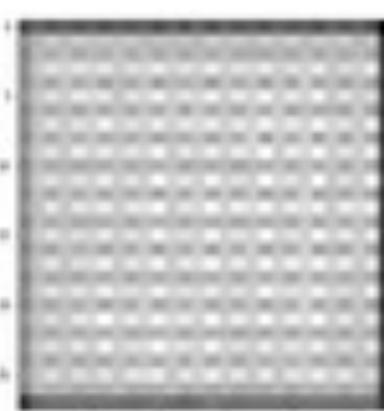
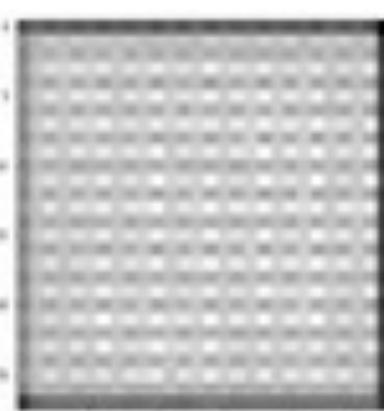
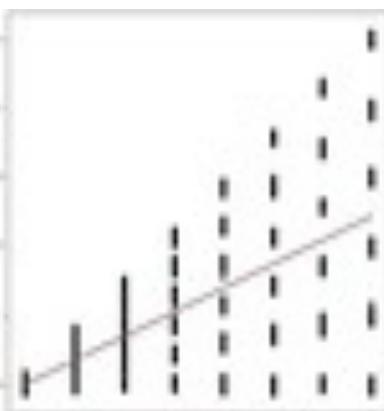
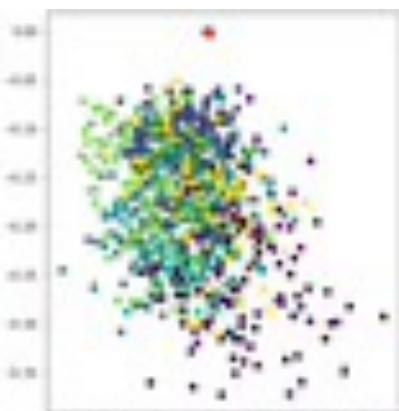
- Gaussian product distribution gives deep posterior embedding for each modality
- Gaussian mixture prior in latent space identifies populations in data across modalities
- *Closed form expressions* for loss – no Monte Carlo
- Supports Bayesian inference across modalities

## Toy code: calibrating a simple affine expert model

[Link to TF  
Code](#)

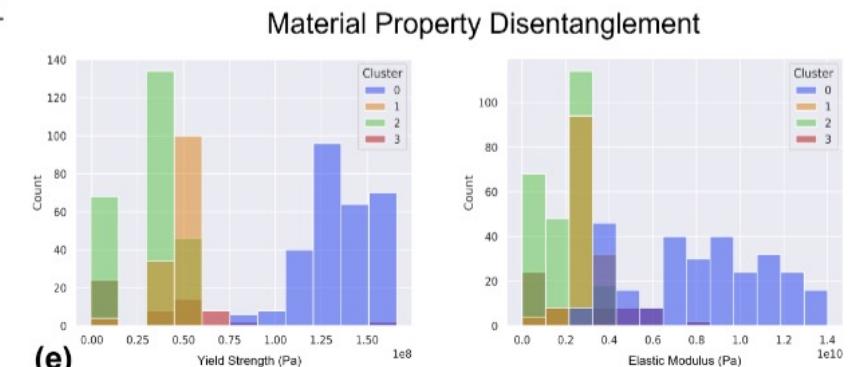
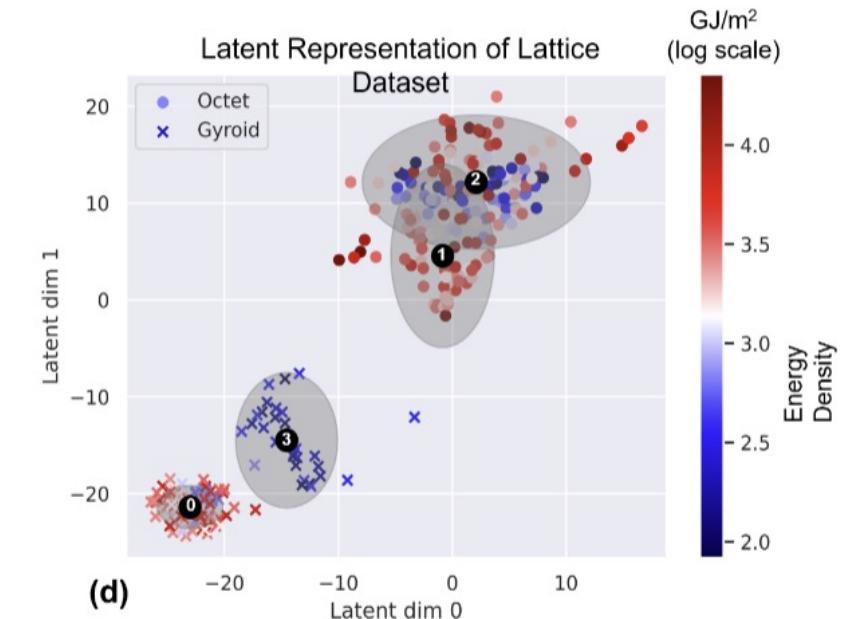
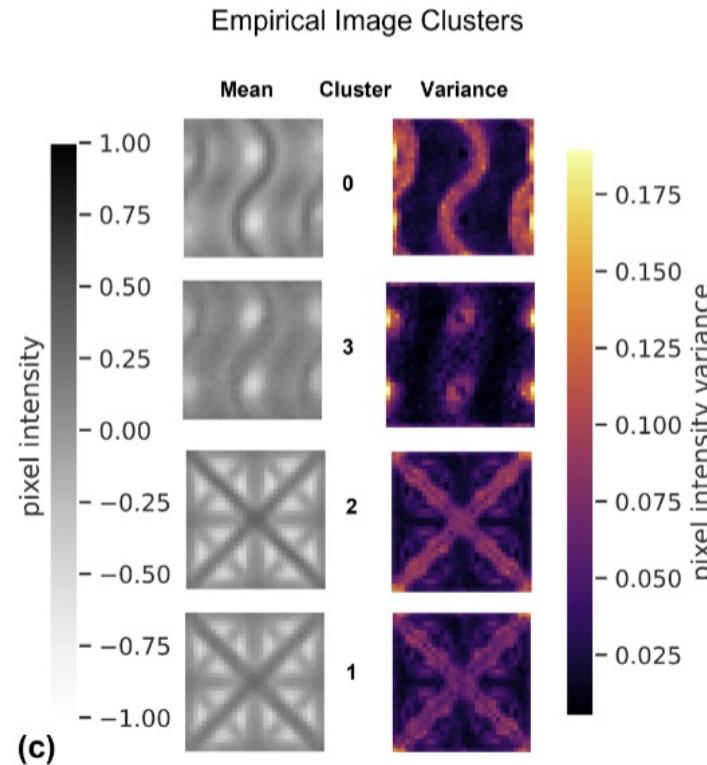
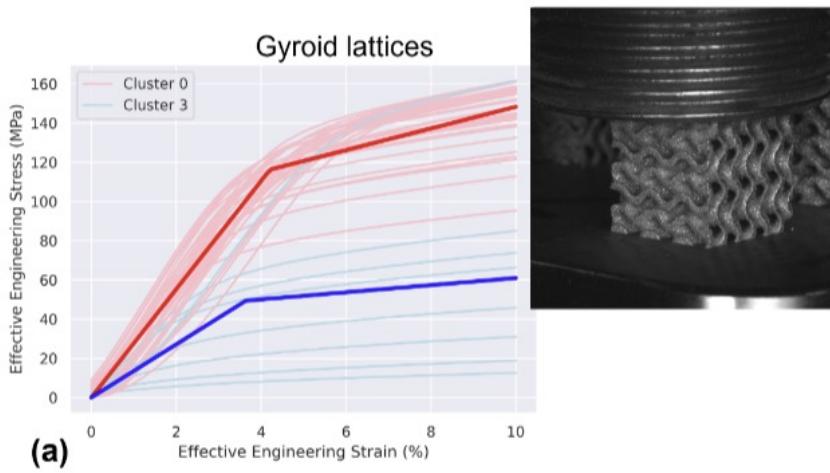
9	8	9	8	8	1	6	8	8	6
8	2	9	2	1	0	1	1	4	2
4	9	1	8	0	5	2	0	4	4
6	0	3	2	0	9	6	2	8	1
8	9	4	1	5	6	1	8	4	9
8	6	4	8	2	9	8	1	5	0
7	2	5	5	5	8	0	9	4	3
9	4	9	5	9	0	9	1	8	1
4	1	4	0	9	1	0	8	3	
1	8	5	0	5	4	2	1	8	7

$X_1$  - Image of Z



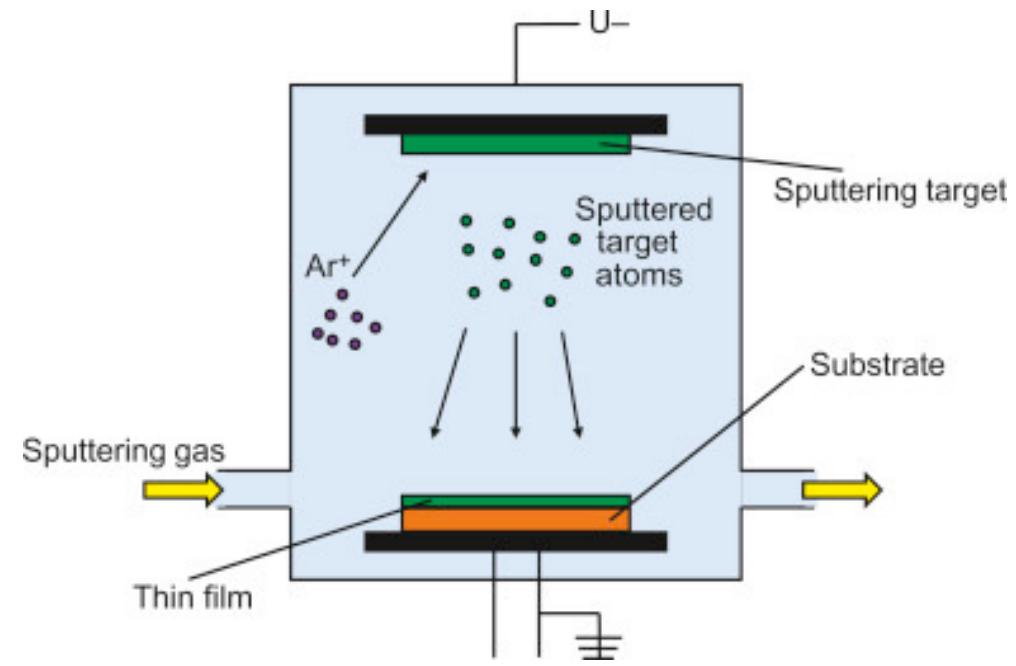
$X_2 = Z^*x + \text{noise}$

# Application 1: Laser powderbed fusion (revisited)

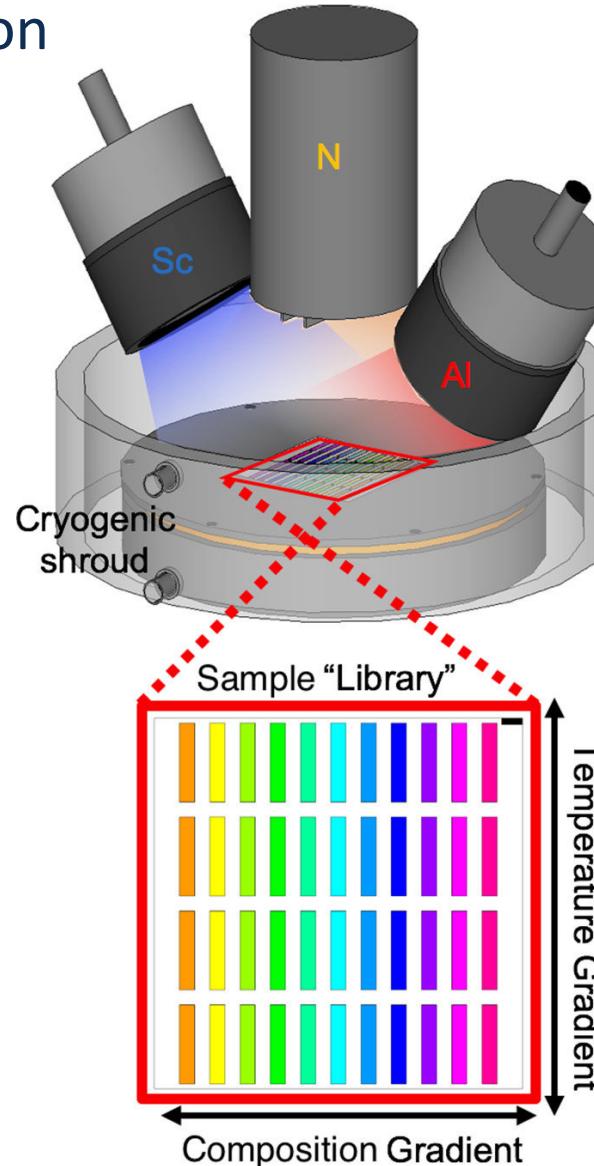


Physical model reveals qualitative regimes of performance,  
other modalities facilitate hypothesis generation for  
explanability

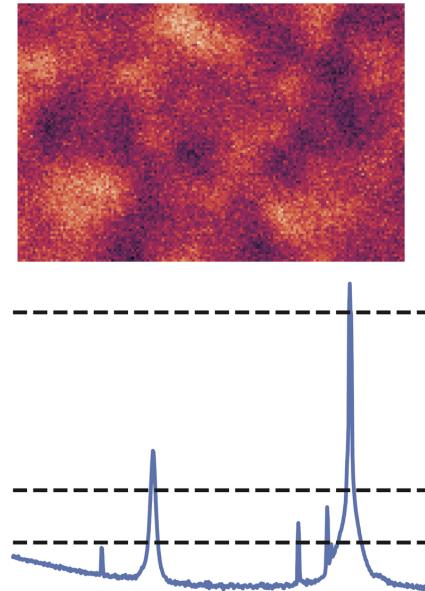
## Application 2: Physical vapor deposition



Standard PVD deposition



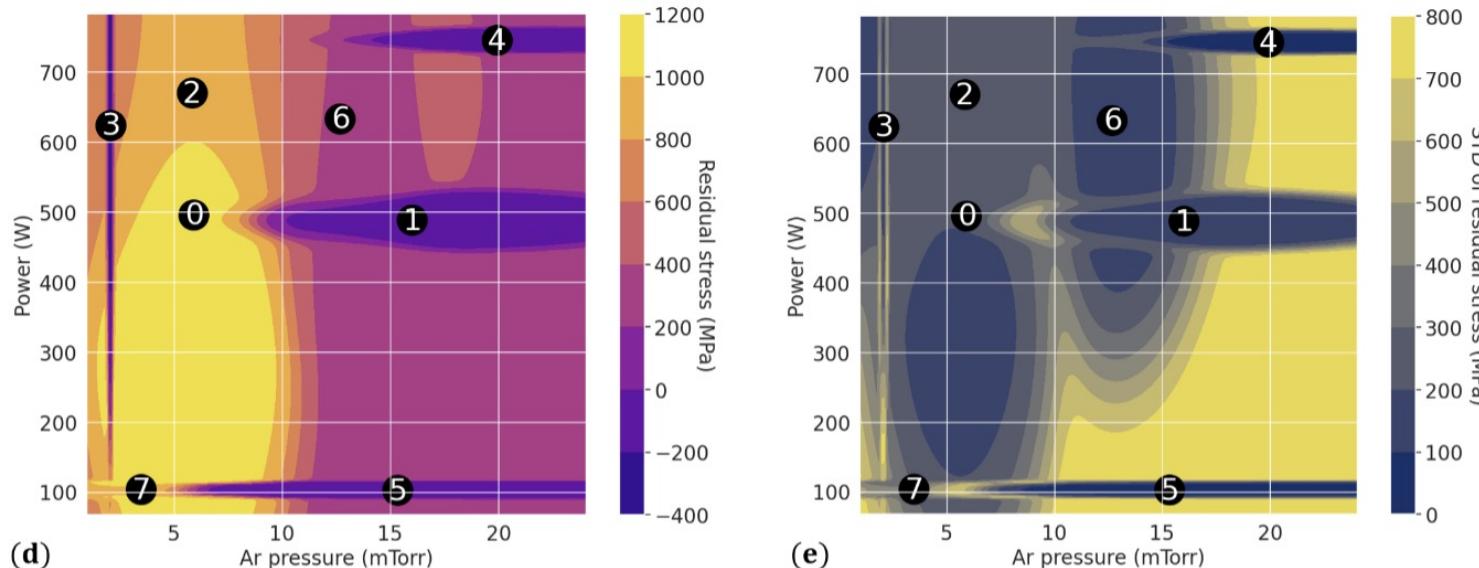
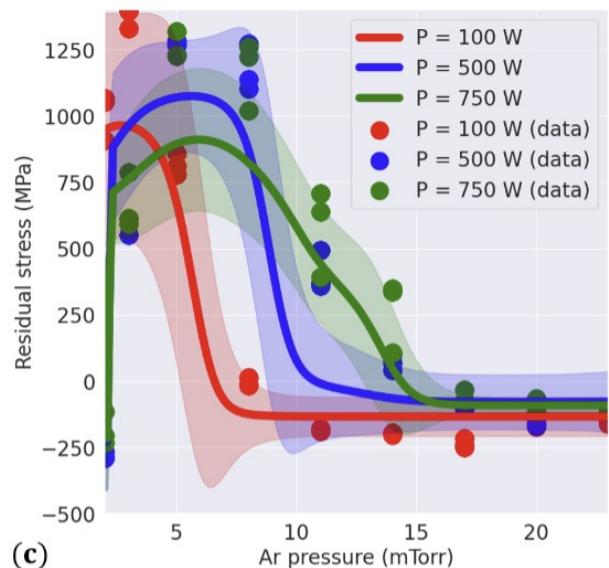
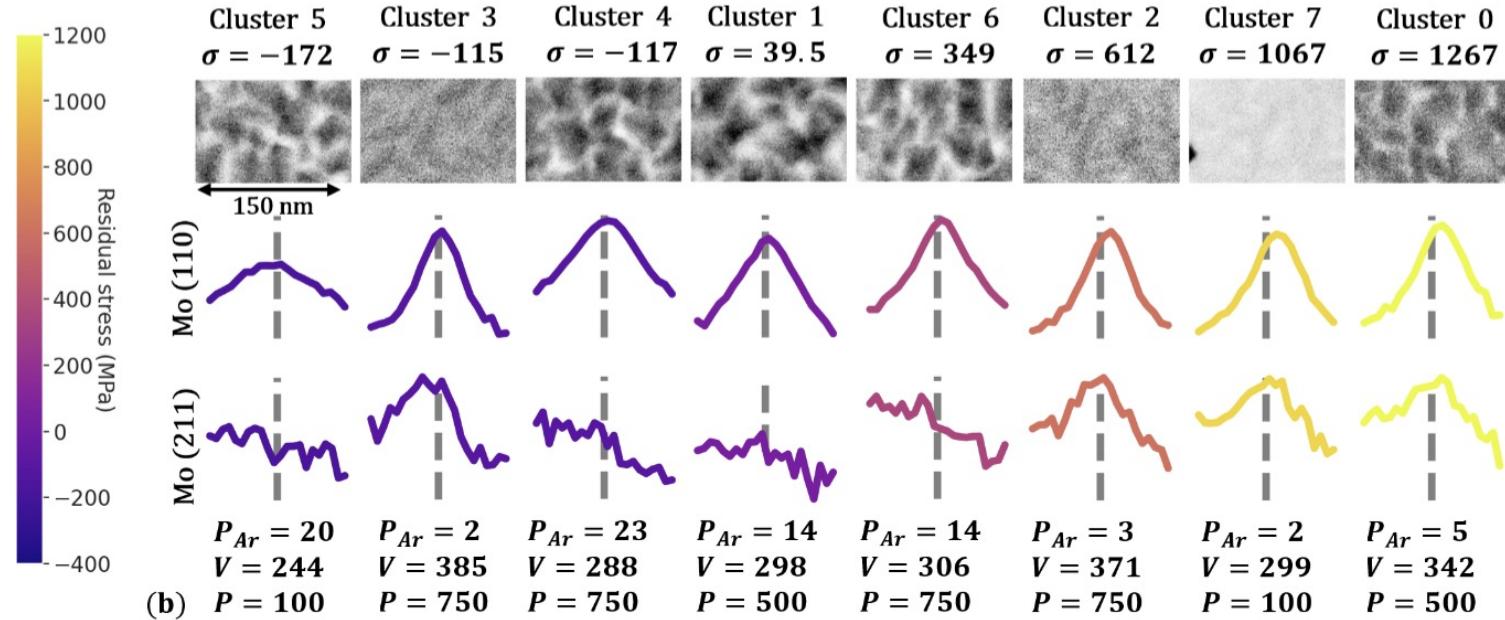
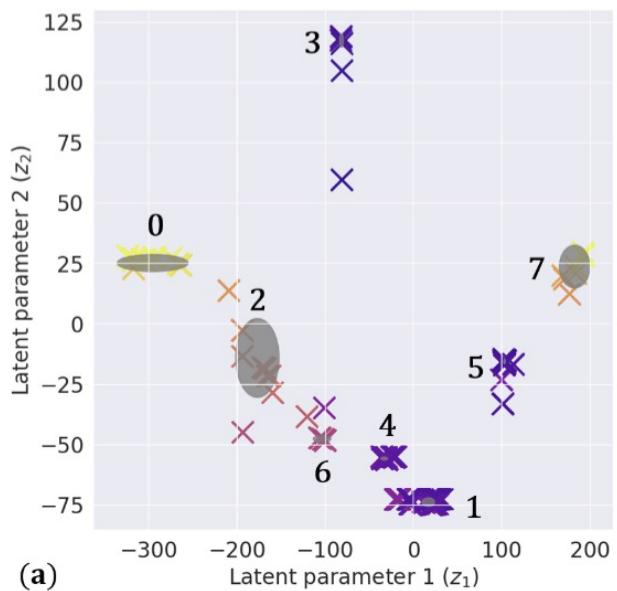
Combinatorial deposition to access large data



### Multimodal characterization

- SEM images
- XRD profiles
- Residual stress measurements
- Process parameters

# Comprehensive multimodal embeddings of physical vapor deposition data



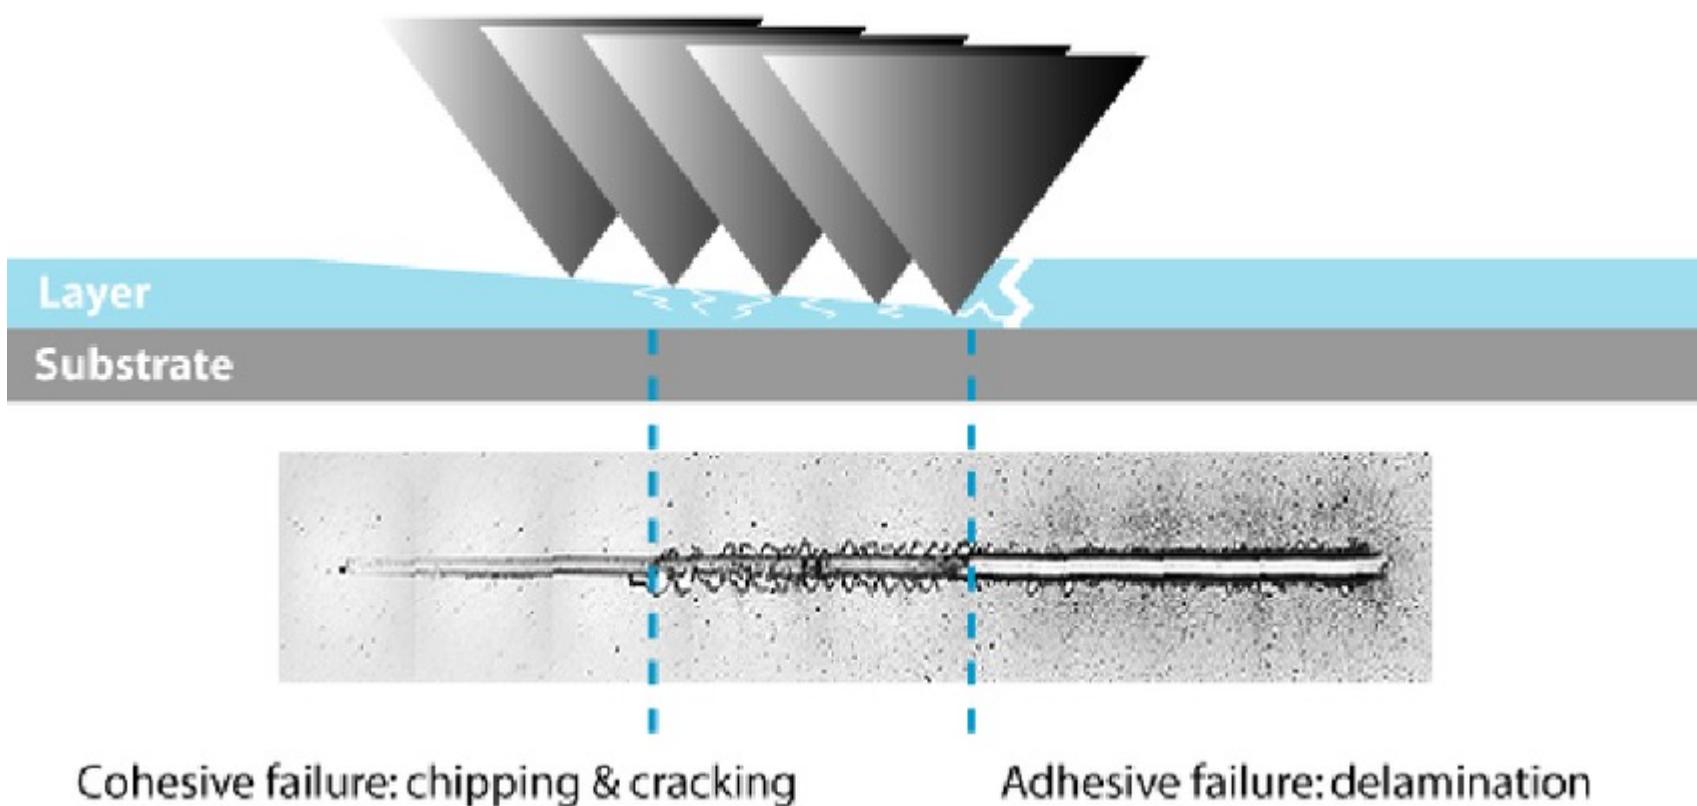
## The real test: are we greater than the sum of the parts?

Input Modalities	Train	Test	Val.
$(X_1, X_2, X_3, X_4)$	1	1	1
$(X_1, X_2, X_3)$	1.225	1.310	1.413
$(X_1, X_2)$	1.827	2.745	3.728
$(X_1, X_3)$	1.286	1.323	1.473
$(X_2, X_3)$	1.286	1.441	1.466
$(X_1)$	2.736	3.528	4.057
$(X_2)$	2.247	2.779	3.618
$(X_3)$	1.640	1.419	1.600

Table 1. Averaged relative  $L_2$  errors for 10 runs for crossmodal prediction of residual stress. Errors are presented as the ratio of the crossmodal prediction w.r.t. using all four modalities (i.e., all four modalities predict residual stress for the train/test/val datasets as 18.67%/23.36%/20.48%). Note:  $X_1$  = SEM images,  $X_2$  = XRD curves,  $X_3$  = process parameters,  $X_4$  = residual stress

## Application 3: tribological testing of platinum gold films

Progressive load measuring depth, friction & acoustic emission



$$p(c|X_{m'_1}, \dots, X_{m'_N})$$

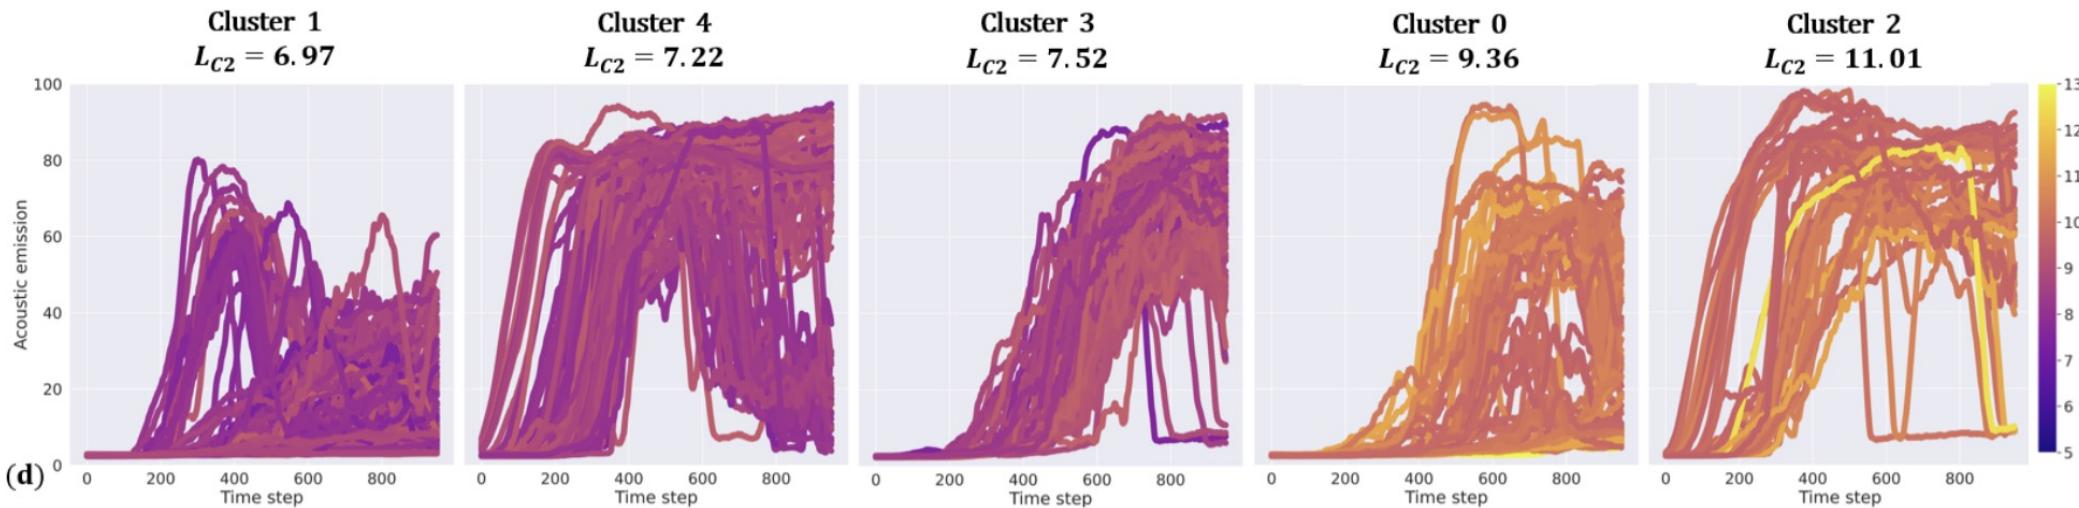
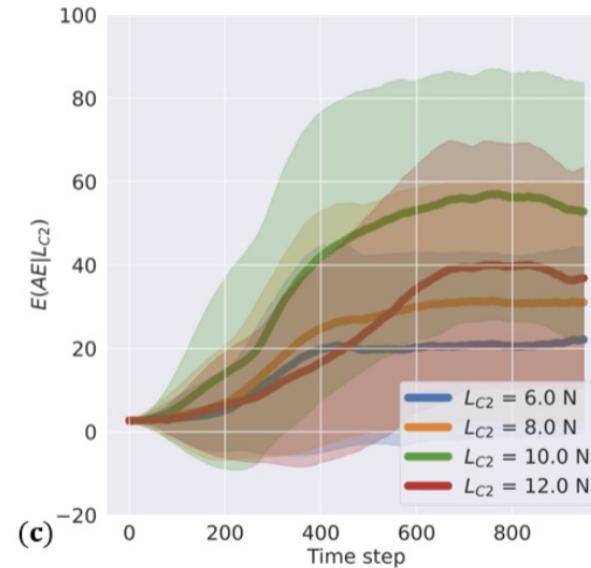
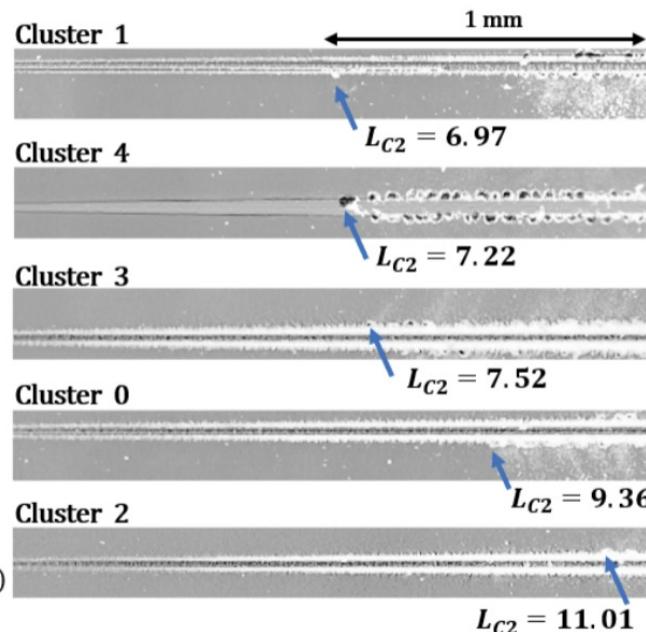
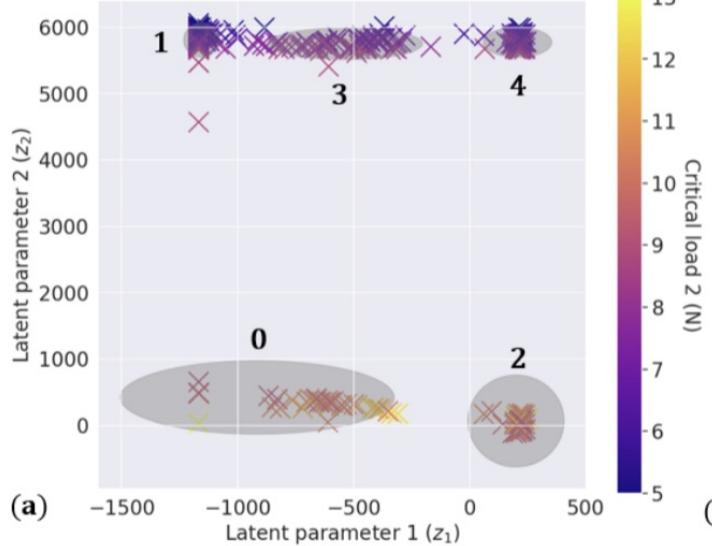
$$= \frac{\prod_{i=1}^N p(X_{m'_i}|c)\pi_c}{\sum_{c'} \prod_{i=1}^N p(X_{m'_i}|c')\pi_{c'}}$$

### Science Question:

Can we use Bayesian estimate of cluster ownership to deduce LC2 from acoustic emission?

If we can, we can build high-throughput experiments to detect unconventional material fingerprints

# Comprehensive multimodal embeddings of physical vapor deposition data



# Closing thoughts

## Embedding physics-informed models alongside physics agnostic data

Not just looking for good compression, but fingerprints which optimally describe relevant physics or support a given task

## How can we use advances in SciML to embed PDE-based expert models?

See next two talks of session

## Toward causality in material fingerprinting

How to identify root causes in fingerprints?

## Modern architectures

Current SotA is data-hungry transformers – are there data-efficient alternatives to variational autoencoders?